

BitVM: Quasi-Turing Complete Computation on Bitcoin

Abstract—A long-standing question in the blockchain community is which class of computations are efficiently expressible in cryptocurrencies with limited scripting languages, such as Bitcoin Script. Such languages expose a reduced trusted computing base, thereby being less prone to hacks and vulnerabilities, but have long been believed to support only limited classes of payments.

In this work, we confute this long-standing belief by showing for the first time that arbitrary computations can be encoded in today’s Bitcoin Script, without introducing any language modification or additional security assumptions, such as trusted hardware, trusted parties, or committees with secure majority. In particular, we present BitVM, a two-party protocol realizing a generic virtual machine by a combination of cryptographic and incentive mechanisms. We conduct a formal analysis of BitVM, characterizing its functionality, system assumptions, and security properties. We further demonstrate the practicality of our approach: in the optimistic case (i.e., in the absence of disputes between parties), our protocol requires just three on-chain transactions, whereas in the pessimistic case, the number of transactions grows logarithmically with the size of the virtual machine. This work not only solves a long-standing theoretical problem, but it also promises a strong practical impact, enabling the development of complex applications in Bitcoin.

1. Introduction

Smart contracts play a fundamental role in blockchain ecosystems by enabling decentralized, automated execution of agreements without the need for trusted intermediaries. These self-executing contracts offer many advantages, such as transparency, security, and immutability. In particular, smart contracts enable programmable money and complex decentralized applications (dApps), fostering innovation in fields like decentralized finance (DeFi), governance, supply chain management, and more.

Smart contracts are typically stored and executed on the blockchain in a low-level language. Some blockchains support a very limited scripting language, such as Bitcoin Script, whereas others feature a quasi-Turing complete¹ language, such as EVM bytecode. The former choice is motivated by a reduced trusted code base, which in principle reduces the attack surface, whereas the latter is justified by

the need to support advanced computations, such as those at the core of DeFi.

A long-standing question within the community is the extent of computational capabilities in Bitcoin Script. This is not only a compelling theoretical question but one with substantial practical implications: if limited scripting languages could support complex computations, they could pave the way for advanced dApps and DeFi applications on secure platforms like Bitcoin. Until now, the prevailing view has been that limited scripting languages, such as Bitcoin Script, are suited primarily for basic functionalities like conditional payments or hashed timelock contracts, with arbitrary computations considered out of reach. Realizing such functionalities would require (i) covenant opcodes [2], which are not yet available in Bitcoin or many other cryptocurrencies, and (ii) costly encoding methods, such as counter machines [3].

Contributions. In this work, we challenge this long-standing belief by showing arbitrary (bounded) computations can be executed securely on Bitcoin in a practical manner. We introduce BitVM, a two-party protocol where a prover P claims that a function evaluates to a specific output for a given input, with a verifier V able to prove fraud and penalize P if the claim is false.

With this mechanism, one can encode any computable function on Bitcoin and execute transactions depending on the function output. For instance, imagine a party V wants to challenge a party P to solve a chess puzzle², say a checkmate-in-1 puzzle, and bet 10 coins (5 coins each) on whether P will solve it in a given timeframe. Note that V may or may not know the solution itself. Both parties lock up the funds in a BitVM instance that, given the initial chessboard configuration and a sequence of chess moves, verifies whether P solves the puzzle in time and distributes the funds accordingly, e.g., if P solves the puzzle in time then 10 coins go to P , otherwise 10 coins go to V .

BitVM does not require any consensus change on Bitcoin and is practical. When parties agree, the on-chain footprint is minimal, needing only three transactions to complete the protocol. In case of a dispute, BitVM ensures a constant upper bound on the number of on-chain transactions.

The key to enabling expressiveness while retaining efficiency is the design of a virtual machine supporting a Turing-complete instruction set, which is then used to execute programs off-chain and produce a verifiable execution

1. This term is adopted in the blockchain community to indicate Turing-complete languages that enforce termination by bounding the execution (e.g., via gas consumption in Ethereum) [1].

2. According to https://en.wikipedia.org/wiki/Chess_puzzle, in a chess puzzle “[...] the goal is to find the single best, ideally aesthetic move or a series of single best moves in a chess position [...]”. We challenge the reader to solve the following checkmate-in-1 puzzles [4], [5].

trace. In case of dispute, the parties can first identify the point of disagreement in the execution trace and then verify a single step of computation on-chain leveraging the Bitcoin Script and the UTXO model, thereby ensuring that disputes are resolved within Bitcoin’s scripting limitations.

The contributions of this paper are summarized below.

- We present BitVM, the first protocol to encode arbitrary computations in Bitcoin Script;
- We conduct a formal analysis of BitVM, characterizing its functionality, system assumptions, and security properties termed balance security and rational correctness;
- We conduct a complexity analysis, detailing on-chain costs and settlement time in both optimistic (no disputes) and dispute scenarios. Specifically, optimistic execution requires only three on-chain transactions, costing approximately 10,707 satoshis³ (6.90\$ in Sept. 2024). In disputes, the on-chain footprint scales logarithmically with the virtual machine’s computational power. For a virtual machine comparable to a high-end ’90s workstation⁴, contract settlement may require up to 81 transactions, costing around 327Ksat, (210\$ in Sept. 2024).

Related work. Relatively complex smart contracts can be implemented in Bitcoin by combining UTXOs and scripts, effectively splitting functionality across multiple transactions. BitML [6] provides a high-level, domain-specific language and compiler that translates programs into Bitcoin transactions, illustrating Bitcoin’s potential for intricate smart contract designs [7]. These methods, however, incur substantial on-chain costs, as compiled programs often result in numerous large transactions that must ultimately be recorded on-chain.

To mitigate these costs, some approaches leverage Trusted Execution Environments (TEEs), such as FastKitten [8] that facilitates off-chain computation within a secure hardware enclave. This approach, however, introduces dependencies on the TEE, a rational adversarial model, collateral, the involvement of a TEE operator, and a limited contract duration. POSE [9] improves upon FastKitten by eliminating the need for collateral and time limitations while enhancing privacy but it still relies on TEE hardware.

Solutions based on Hashed Timelock contracts (HTLCs) aim to shift computation off-chain in a way akin to *state channels* on Ethereum [10], [11] while remaining compatible with Bitcoin’s inherent constraints [12]. For example, *Discreet Log Contracts (DLCs)* [13] and *Cryptographic Oracle-Based Conditional Payments* [14] are alternative approaches that utilize (semi-)trusted oracles to assert specific events and perform payments conditioned to them. These events, typically encoded in the preimage of the hash, have to be known upfront, which restricts the class of supported functions. For instance, the chess example from Section 1 would not be expressible through HTLCs since the outcome might not be known a-priori to any of the parties.

3. A *satoshi* is a fraction of a bitcoin, i.e., $1\text{sat} = 10^{-8}\$$.

4. Configured with 2^{32} memory cells for 32-bit integers, supporting up to 2^{32} instructions and steps.

In contrast to prior work, BitVM enables quasi-Turing complete computation on Bitcoin, similar to Ethereum, without relying on additional assumptions such as TEEs or semi-trusted oracles. BitVM thus represents the first trustless protocol allowing arbitrary, yet bounded, computation on Bitcoin, unlocking a range of potential applications. Notable examples include bridges (e.g., [15], [16], [17]), some of which are already being deployed based on the initial informal BitVM concept [18].

Due to the industry’s significant interest in this concept, a follow-up work proposed an alternative approach with the main goal of designing a bridge between Bitcoin and layer-2 systems [19]: The core idea in that work is to compile programs down to potentially huge Bitcoin Script programs, split them, and commit to intermediary results on-chain, which can then be disputed. This approach allows permissionless challengers to dispute false claims of correctness but incurs high on-chain costs. In case of dispute in [19], the protocol enforces at least one transaction on-chain that fills an entire 4 MB Bitcoin block, costing approximately 1.9K\$. In contrast, the solution we present in this paper is better suited to permissioned environments, as it is significantly more cost-effective in terms of on-chain fees compared to [19]; the estimated cost for BitVM in case of dispute is approximately 210\$.

Approach	Expressiveness	Extra Assumptions	On-chain cost
BitML [6], [7]	QT	None	$O(n)$
TEEs [8], [9]	QT	TEE	$O(n)$
Gen. Channels [12]	Bitcoin	None	$O(1)$
Oracles [13], [14]	QT	trusted oracle	$O(1)$
BitVM	QT	None	$O(\log(n))$

TABLE 1. COMPARISON OF BITCOIN-BASED SMART CONTRACT APPROACHES. n DENOTES THE UPPER BOUND ON COMPUTATIONAL STEPS, AND QT REFERS TO QUASI-TURING COMPLETENESS.

2. Model

BitVM is a protocol between two mutually distrusting parties, the prover P and the verifier V , designed to enable P to prove on the Bitcoin blockchain that the outcome of a pre-agreed computation with V was performed correctly. Concretely, for an agreed-upon Turing-complete program Π , a BitVM instance secures collateral from both parties and it enables P to enforce a transaction on-chain based on the outcome $\Pi(x)$ for a specific input x . In other words, $\Pi(x)$ dictates the payout of the funds within the BitVM instance, typically allocating them to P and V ⁵. If P or V stop collaborating during protocol execution, after a designated period all the funds are allocated to the other party.

2.1. System model

We assume time advances in discrete rounds $(1, 2, \dots)$. Protocol participants run in probabilistic polynomial time

5. Note that P and V can agree to allocate the funds to a third party or, more generally, make the funds spendable under any condition that can be expressed in Bitcoin Script.

(PPT) in the security parameter κ . We assume synchronous communication, i.e., messages sent between parties arrive at the beginning of the next round, as well as authenticated communication channels. Our protocol employs a hash function modeled as a random oracle $H : \{0, 1\}^* \rightarrow \{0, 1\}^\kappa$ which maps an input of arbitrary length to a fixed κ -sized output. Moreover, our protocol builds upon a distributed ledger protocol (e.g., [20], [21], [22]).

Definition 1 (Distributed Ledger Protocol). A distributed ledger protocol is an interactive Turing machine exposing the following functionality on each party.

- `execute()`: executes one protocol round and enables the machine to communicate with the network, invoked by the environment in every round;
- `write(tx)`: takes as input a transaction from the environment;
- `read()`: outputs a finite, ordered sequence of transactions, also known as *transaction ledger* L .

We denote L_r^P as the output of invoking `read()` on party P at the end of round r . We restrict honest parties to only include *valid* transactions in their ledgers⁶. As we are interested in building BitVM on Bitcoin, when we present the construction, transactions are deemed (in)valid based on Bitcoin’s validation rules (see Section 3.1). However, BitVM can be built on top of any distributed ledger protocol with validation rules as expressive as those of Bitcoin. We assume that our protocol participants have access to the functionality exposed by the distributed ledger protocol, either by being an active participant or by running some (light) client protocol. We are interested in distributed ledger protocols that are safe and live, as defined below (cf. [20], [21], [22]). Given two sequences A and B , we use $A \preceq B$ to mean that A is a prefix of B .

Definition 2 (Stickiness). A distributed ledger protocol is sticky if for any honest party P and any rounds $r_1 \leq r_2$, it holds that $L_{r_1}^P \preceq L_{r_2}^P$.

Definition 3 (Safety). A distributed ledger protocol is safe, if it is sticky and for any pair of honest parties P_1, P_2 and any pair of rounds r_1, r_2 , it holds that $L_{r_1}^{P_1} \preceq L_{r_2}^{P_2} \vee L_{r_2}^{P_2} \preceq L_{r_1}^{P_1}$.

Definition 4 (Liveness). A distributed ledger protocol execution is live(u), if any transaction that is written to an honest party’s ledger at round r , appears in the ledger of all honest parties by round $r + u$, denoted as L_{r+u}^\cap .

Throughout this paper, we say “publish a transaction tx (on L)” to denote calling the function `write(tx)`. Furthermore, after publishing a valid transaction tx , we sometimes say “wait until tx appears (on L)”, to denote calling the function `read()` every round until $tx \in L$, which happens at most after u rounds due to liveness. When presenting the BitVM construction, we sometimes refer to the ledger as blockchain even though the distributed ledger protocol could be realized

differently. We say something happens *on-chain* if there are one or more corresponding transactions in the ledger, and something happens *off-chain* if there are no corresponding transactions on the ledger.

There is a ledger state that is induced by a ledger L , denoted as $\text{st}(L)$, by executing each transaction in order, starting with a genesis state. The execution of transactions is captured by a state transition function, taking a state and a transaction and outputting a new state. We denote $\text{bal}_L(P) \in \mathbb{R}_{\geq 0}$ as the balance of party P in the state induced by L . A party can use parts of their balance in $p \in [0, \text{bal}_L(P)]$ as *monetary input* for a transaction. For a given ledger L , we define the on-chain (monetary) utility of a transaction $tx \in L$ for a party P as $w_L(P, tx) := \text{bal}_{L_1}(P) - \text{bal}_{L_2}(P)$, where $L_1 \prec L$ is the ledger up to (not including) tx and $L_2 := L_1 || tx$. Usually, it is obvious which ledger we refer to, so we omit the subscript. In addition to balances of parties, a ledger state $\text{st}(L)$ can include a string $s \in \{0, 1\}^*$, denoted as $s \in \text{st}(L)$, if there exists a transaction $tx \in L$, such that tx contains the string s .

2.2. Threat model

We analyze BitVM in the presence of a PPT adversary that may corrupt any protocol party $\{P, V\}$ during the execution of the protocol. The adversary can corrupt parties, causing them to behave either as *Byzantine* or as *rational* actors. Byzantine parties can deviate arbitrarily from the honest protocol execution. Contrarily, rational parties deviate from the honest protocol execution only when such action increases their monetary utility.

The protocol gives different guarantees based on the type of corruption. On a high level, we want to show that (i) honest protocol participants are guaranteed their rightful balance even if the other party is Byzantine, (ii) rational parties follow the honest protocol execution, and (iii) if both parties behave rationally, the protocol follows an optimistic execution (which is efficient). We formally define these properties in Section 2.3.

2.3. Protocol goals

The core objectives of BitVM are termed *balance security* and *rational correctness*. Informally, balance security ensures an honest party will not lose their funds against Byzantine counterparties, whereas rational correctness guarantees that rational parties will follow the protocol. To formally define balance security we argue in terms of utility, i.e., the utility of the on-chain state of an honest party after the settlement of a BitVM instance will be at least equal to its utility of the correct final state, regardless of the actions of its counterparty. Rational correctness implies that if both parties are rational, they will commit on-chain the correct final state of the BitVM instance. These properties are standard in the literature: for instance, an honest user of a Lightning channel [23] can always dispute a malicious commitment and claim the channel funds, while rational players will always commit to the last agreed-upon state [24].

⁶. This is not strictly necessary and is done mainly for convenience. Parties could also take an outputted ledger and remove invalid transactions from it.

We formalize these objectives on a generic primitive, which we call *on-chain state verification* protocol and is defined as follows.

Definition 5 (On-chain State Verification Protocol). An on-chain state verification protocol, parameterized over a distributed ledger protocol that outputs a ledger L , is a two-party protocol that exposes the two following functionalities:

- $\text{setup}(\text{in}_P, \text{in}_V, \Pi, f)$: takes as input monetary inputs $\text{in}_P \in [0, \text{bal}_L(P)]$ and $\text{in}_V \in [0, \text{bal}_L(V)]$ of parties P and V , a computable function (or program) $\Pi : \mathcal{S} \rightarrow \mathcal{O}$ that maps a set of states \mathcal{S} to a set of outcomes \mathcal{O} and an outcome mapping function $f : \mathcal{O} \rightarrow \mathbb{R}_{\geq 0}^2$, that maps the set of outcomes \mathcal{O} to pairs of utilities (v_P, v_V) where $v_P + v_V \leq \text{in}_P + \text{in}_V$ and returns an instance \mathcal{I} .
- $\text{execute}(\mathcal{I}, x)$: takes as input an instance \mathcal{I} returned by the setup function and a function input $x \in \mathcal{S}$ (for function Π).

Consider an execution of this primitive for given inputs $\text{in}_P, \text{in}_V, \Pi, f$, where $\mathcal{I} \leftarrow \text{setup}(\text{in}_P, \text{in}_V, \Pi, f)$, and then $\text{execute}(\mathcal{I}, x)$ are called, and finish in round r . Let \mathcal{T} be the set of transactions that are included in L_{r+u}^\cap as a result of this execution. Moreover, we denote the utility of party $A \in \{P, V\}$ in $f(\Pi(x))$ by $f_A(\Pi(x))$.

Balance Security. An execution achieves balance security, if it holds that $\sum_{tx \in \mathcal{T}} (w(tx, A)) \geq v_A$ where $v_A = f_A(\Pi(x))$, for any honest $A \in \{P, V\}$.

Rational Correctness. An execution achieves rational correctness, if P and V are rational and $\sum_{tx \in \mathcal{T}} (w(tx, A)) = v_A$ where $v_A = f_A(\Pi(x))$, for any $A \in \{P, V\}$ and $\Pi(x) \in \text{st}(L_{r+u}^\cap)$.

An on-chain state verification protocol achieves balance security and rational correctness, respectively, if for any $\text{in}_P, \text{in}_V, \Pi, f$ the probability that the corresponding execution does not achieve balance security and rational correctness, respectively, is negligible in κ .

3. Preliminaries

In this section, we present the necessary background concerning Bitcoin Script and some key primitives our construction builds upon.

Notation. Given a sequence $A := (a_1, \dots, a_n)$, $A[i]$ represents its i -th element. We use $A[i : j]$ to denote the subsequence (a_i, \dots, a_j) . We use $|A|$ to denote the length of a sequence, e.g., $|(a_1, \dots, a_n)| = n$. For a string $s \in \{0, 1\}^*$, we use $|s|_{\text{bit}}$ to denote its bit length.

3.1. Transactions in the UTXO model

A user U on a ledger L is identified by the secret-public key pair $(\text{pk}_U, \text{sk}_U)$; by $\sigma_U(m)$ we denote the digital signature of U over the message $m \in \{0, 1\}^*$.

In the *unspent transaction output* (UTXO) model, a transaction Tx maps a (non-empty) list of existing, unspent, transaction outputs to a (non-empty) list of new transaction

outputs. A transaction output is defined as an attribute tuple $\text{out} := (a\$, \text{lockScript})$, where $\text{out}.a \in \mathbb{R}_{\geq 0}$ is the amount of coins (expressed in $\$$) held by the output out and out.lockScript is the condition that needs to be fulfilled to spend it and transfer the coins to a new output, which we also call UTXO. We distinguish the already existing transaction outputs (input of a transaction Tx) from the newly created outputs calling them Tx.inputs and Tx.outputs , respectively. A transaction input in is defined as $\text{in} := (\text{PrevTx}, \text{outIndex}, \text{lockScript})$, where the output being spent is uniquely identified by specifying the transaction PrevTx and an output index outIndex . To improve readability, we also give the locking script lockScript that is being fulfilled.

We formally define a transaction as a tuple $\text{Tx} := (\text{inputs}, \text{witnesses}, \text{outputs})$ where $\text{Tx.inputs} := [\text{in}_1, \dots, \text{in}_n]$ are the transaction inputs, $\text{Tx.outputs} := [\text{out}_1, \dots, \text{out}_m]$ are the transaction outputs and $\text{Tx.witnesses} := [w_1, \dots, w_n]$ represents the witness data, i.e., the list of the tuples that fulfill the spending conditions of the inputs, one witness for each input. The locking script of an output is expressed in the scripting language of the ledger. To transfer the coins held in a UTXO, its locking script is executed with a witness as script input and must return `True`; if successful, the condition is considered fulfilled. If the script execution returns `False`, the condition is not fulfilled and the UTXO is not spendable⁷.

A transaction is valid only if every UTXO in input is unspent, the witnesses fulfill the conditions of the corresponding locking scripts, and the sum of the coins held in the inputs is equal to or greater than the sum of the coins held in the outputs.

Transaction spending conditions. Bitcoin has a stack-based scripting language. Below, we describe the subset of Bitcoin spending conditions that we use in this paper.

- **Signature locks.** The spending condition $\text{CheckSig}_{\text{pk}_U}(m)$ is fulfilled if the signature $\sigma_U(m)$ is part of the witness.
- **Multisignature locks.** To fulfill this spending condition, k out of n signatures are required. In particular, for two users A and B , a spending condition that represents a 2-of-2 multi-signature of a message m between them is denoted as $\text{CheckMSig}_{\text{pk}_{A,B}}(m)$ and is fulfilled by giving the signature $\sigma_{A,B}(m)$ as part of the witness of the spending transaction.
- **Relative timelocks** make a transaction output spendable only after a specified time Δ has elapsed since the transaction was included on-chain. We denote the relative timelock spending condition as $\text{TL}(\Delta)$.
- **Taproot trees** [26], also known as Taptrees, enable a UTXO to be spent by satisfying one of several possible spending conditions. These conditions, referred to as

7. In this work, we separate the locking script from the witness for readability. However, note that in practice, the protocol is implemented using SegWit [25] transactions, where the locking script is included in the witness.

Tapleaves, form the leaves of a Merkle tree. To spend a UTXO locked by a Taptree locking script, the user must provide a witness for one of the Tapleaves along with proof of inclusion of that leaf in the Taptree.

We denote the Tapleaves of a Taptree locking script as $\langle \text{leaf}_1, \dots, \text{leaf}_r \rangle$. When a user fulfills the script leaf_i to unlock the j -th output of the transaction Tx , the corresponding input is represented as $(\text{Tx}, j, \langle \text{leaf}_i \rangle)$.

Whenever a user spends a UTXO via a Tapleaf of a Taptree, we assume that they have provided a valid Merkle proof of inclusion for that Tapleaf.

- **Other conditions.** We denote with True (False) a condition that is always fulfilled (can never be fulfilled).

We use $*$ to denote a generic transaction input, witness, or output that is not directly relevant to our protocol, provided it remains valid under Bitcoin consensus rules.

Combining spending conditions. When presenting spending conditions with complex logic, we explicitly provide their pseudocode. We use the conditions described in this section as building blocks, combining them with standard Bitcoin Script constructions using logical operators \wedge (and) and \vee (or). Furthermore, for convenience, inside long scripts we append the keyword `Verify` to sub-spending conditions that return either True or False with the following meaning: if the sub-spending condition returns True, pop True from the stack and continue to execute the rest of the script, if it returns False, mark the transaction as invalid (and thus fail to unlock the long script). This is meant to mimic how the Bitcoin `OP_VERIFY` opcode works.

3.2. Lamport digital signature scheme

Let $h : X \rightarrow Y$ be a one-way function, where $X := \{0, 1\}^*$ and $Y := \{0, 1\}^\lambda$, for a given security parameter λ . Let $m \in \{0, 1\}^\ell$ be a ℓ -bit message, with $\ell \in \mathbb{N}_{>0}$. A *Lamport digital signature scheme* [27] Lamport consists of a triple of algorithms (KeyGen, Sig, Vrfy), where:

- $(pk_{\mathcal{M}}, sk_{\mathcal{M}}) \leftarrow \text{Lamp.KeyGen}(\ell)$ (cf. Algorithm 1), is a Probabilistic Polynomial Time (PPT) algorithm that takes as input a positive integer ℓ and returns a key pair, consisting of a secret key $sk_{\mathcal{M}}$ and a public key $pk_{\mathcal{M}}$ which can be used for one-time signing an ℓ -bit message. We use $\mathcal{M} = \{0, 1\}^\ell$ as an alias for the ℓ -bit message space.
- $c_m \leftarrow \text{Lamp.Sig}_{sk_{\mathcal{M}}}(m)$ (cf. Algorithm 2), is a Deterministic Polynomial Time (DPT) algorithm parameterized by a secret key $sk_{\mathcal{M}}$, that takes as input a message $m \in \mathcal{M}$ and outputs the signature c_m , which we also call (Lamport) commitment.
- $\{\text{True}, \text{False}\} \leftarrow \text{Lamp.Vrfy}_{pk_{\mathcal{M}}}(m, c_m)$ (cf. Algorithm 3), is a DPT algorithm parameterized by a public key $pk_{\mathcal{M}}$ that takes as input a message m , a signature c_m , and outputs True iff c_m is a valid signature for m generated by the secret key $sk_{\mathcal{M}}$, corresponding to $pk_{\mathcal{M}}$, i.e., $(pk_{\mathcal{M}}, sk_{\mathcal{M}})$ is a key pair generated by `Lamp.KeyGen`.

Lamport signatures are secure one-time signatures. Given a message space \mathcal{M} , it is possible to sign any message

Algorithm 1 The key generation algorithm `Lamp.KeyGen` for a ℓ -bit messages space \mathcal{M} . In the following algorithms, we use matrix notation, i.e., for a given two-dimensional matrix a , $a[i, j]$ refers to the element at row i and column j of it.

```

1: function Lamp.KeyGen( $\ell$ )
2:   Let  $sk_{\mathcal{M}} \leftarrow \begin{pmatrix} x[0, 0], \dots, x[0, \ell - 1] \\ x[1, 0], \dots, x[1, \ell - 1] \end{pmatrix}$ , where every element  $x[i, j]$  is sampled uniformly at random from the set  $X$ ;
3:   for  $i = 0, 1$  and  $j = 0, \dots, \ell - 1$  do
4:      $y[i, j] \leftarrow h(x[i, j])$ ;
5:   Let  $pk_{\mathcal{M}} \leftarrow \begin{pmatrix} y[0, 0], \dots, y[0, \ell - 1] \\ y[1, 0], \dots, y[1, \ell - 1] \end{pmatrix}$ ;
6:   return  $(sk_{\mathcal{M}}, pk_{\mathcal{M}})$ .
```

Algorithm 2 The Lamport signature algorithm `Lamp.Sig`, parameterized over a secret key $sk_{\mathcal{M}}$ for a ℓ -bit sized message space \mathcal{M} .

```

1: function LampSig $_{sk_{\mathcal{M}}}(m)$ 
2:   for  $i = 0, \dots, \ell - 1$  do
3:     Let  $c_m[i] \leftarrow sk_{\mathcal{M}}[m[i], i]$ ;
4:   return  $c_m$ .
```

$m \in \mathcal{M}$ by using the secret key $sk_{\mathcal{M}}$ of the key pair $(sk_{\mathcal{M}}, pk_{\mathcal{M}})$, i.e., the key pair associated to \mathcal{M} . When the message m is signed and c_m is created, the key pair becomes bound to m . No polynomially bounded adversary is able to forge a signature for a different message $m' \neq m$ with non-negligible probability. However, if the signer uses the same secret key $sk_{\mathcal{M}}$ to sign another different ℓ -bit messages $m'' \neq m$, they can be held accountable. We call this action *equivocation* and we show how to detect it in Algorithm 4.

Notice that signing the ℓ -bit message m with the secret key $sk_{\mathcal{M}}$ consists in revealing for every bit $i = 0, \dots, \ell - 1$ of m one of the two preimages that compose the i -th column of secret key $sk_{\mathcal{M}}$, namely, revealing $x[0, i]$ to claim that $m[i] = 0$, or revealing $x[1, i]$ to claim that $m[i] = 1$. When the signer reveals both $x[0, i], x[1, i]$ for any bit i , they are equivocating.

For a formal discussion about one-time security and a proof that Lamport signatures are one-time secure (assuming the existence of one-way functions), see, e.g., [28]. One-time security is crucial for the correctness of BitVM as it enables the signer of a message to make a non-repudiable commitment to that message. Lamport signatures are implementable using Bitcoin Script, as demonstrated in [29].

In the following, we are interested in Lamport signatures as a mechanism to enable a party to *commit* to (single or multiple bits) messages. Thus, we will refer to Algorithm 2

Algorithm 3 Lamport verification algorithm `Lamp.Vrfy`, parameterized over a public key $pk_{\mathcal{M}}$ for a ℓ -bit message space \mathcal{M} .

```

1: function Lamp.Vrfy $_{pk_{\mathcal{M}}}(m, c_m)$ 
2:   for  $i = 0, \dots, \ell - 1$  do
3:     if  $h(c_m[i]) \neq pk_{\mathcal{M}}[m[i], i]$  then
4:       return False;
5:   return True.
```

Algorithm 4 The CheckEquivocation algorithm for a bit $b \in \mathcal{B} = \{0, 1\}$. The input is the corresponding public key $pk_{\mathcal{B}}$ and two preimages $x', x'' \in X$.

```

1: function CheckEquivocation( $pk_{\mathcal{B}}, x', x''$ )
2:   if ( $h(x') = pk_{\mathcal{B}}[0, 0]$  and  $h(x'') = pk_{\mathcal{B}}[1, 0]$ ) then
3:     return True;
4:    $\triangleright$  The committer is trying to commit to both 0 and 1
      for the bit  $b$ .  $\triangleleft$ 
5:   else
6:     return False.

```

as Comm instead of Lamp.Sig and to Algorithm 3 as CheckComm instead of Lamp.Vrfy.

3.3. Stateful Bitcoin scripting

Although the Bitcoin scripting language is stateless, a clever use of one-time digital signature schemes, such as Lamport signatures, enables state preservation across different Bitcoin transactions.

Consider the following example: Let a user U hold a Lamport key pair $(sk_{\mathcal{M}}, pk_{\mathcal{M}})$ associated with \mathcal{M} , the set of all ℓ -bit messages. We can think of \mathcal{M} as a variable that can hold any ℓ -bit string. U can assign a value m to \mathcal{M} by creating the commitment $c_m \leftarrow \text{Comm}_{sk_{\mathcal{M}}}(m)$.

By hard-coding $\text{CheckComm}_{pk_{\mathcal{M}}}$ for a public key $pk_{\mathcal{M}}$ in the locking script of multiple outputs, this variable assignment can not only be verified but also transferred from one output to another, effectively establishing a global state in Bitcoin. This is accomplished by reading m and c_m from the unlocking script of one output and passing them to another output through its witness. For example, consider two different transactions $\text{Tx}_1 := (*, *, [\text{out}_1, *])$ and $\text{Tx}_2 := (*, *, [\text{out}'_1, *])$, where the outputs are defined as $\text{out}_1 := (a\$, \text{CheckComm}_{pk_{\mathcal{M}}})$ and $\text{out}'_1 := (b\$, \text{CheckComm}_{pk_{\mathcal{M}}})$. To unlock both out_1 and out'_1 , a Lamport commitment c_m must be provided. Since the same Lamport public key appears in both scripts, every party in the network knows that when U unlocks these scripts, U is assigning a value to the same variable \mathcal{M} . Following from *one-time security*, no user other than U can assign a different value to \mathcal{M} without knowing $sk_{\mathcal{M}}$. Moreover, U cannot assign two different values $m_1 \neq m_2$ to \mathcal{M} without equivocating, which is detectable and can be punished on-chain.

4. BitVM Virtual Machine

In the BitVM protocol, both parties employ a *Virtual Machine* (VM) to run off-chain any deterministic program Π . Although the underlying concept closely resembles an *abstract machine*, we choose to retain the term “VM” to stay consistent with the original naming of the construction. In this section, we describe the components of the VM and demonstrate how to initialize them for practical deployment of the protocol.

VM components. At a high level, the virtual machine (VM) executes programs composed of instructions written

in a VM-compatible language. While the program is running, the VM continuously performs an instruction cycle, or *state transition function*. In each cycle, the VM fetches the instruction indicated by the program counter, loads the values stored at specific memory addresses referenced by the instruction, executes the operation defined by the instruction on those values, stores the result at the designated memory address, and updates the program counter accordingly (cf. Definition 7).

This process repeats until the program terminates or reaches a predefined execution limit. Throughout its execution, the VM produces an execution trace, recording (i) the current program counter value and (ii) a commitment to the state of memory at each step. The BitVM protocol leverages this execution trace for dispute resolution, as described in Section 6.3 and Section A.1.

Formally, let a *VM address* be an integer $addr \in \mathcal{A} := \{0, 1, \dots, \text{MemLen} - 1\}$ where $\text{MemLen} \in \mathbb{N}_{>0}$ represents the memory length. We define the *VM memory* as the sequence $M \in \mathcal{M} := \{0, 1, \dots, n\}^{\text{MemLen}}$, where $n \in \mathbb{N}_{>0}$ specifies the range of values stored at any memory address. The *VM program counter*, denoted pc , is an element of the set $\mathcal{PC} := \{0, 1, \dots, \ell - 1\} \cup \{\perp\}$, where $\ell \in \{1, 2, \dots, n\}$ is the maximum length of the program, and \perp indicates termination. Let $\mathcal{OP} := \{f_{\text{OP}} : \mathcal{PC} \times \{0, \dots, n\} \times \{0, \dots, n\} \rightarrow \mathcal{PC} \times \{0, \dots, n\} \cup \{\perp\}\}$ be a set of CPU instructions that the VM can execute⁸. The function f_{OP} takes as input a triple (pc, val_A, val_B) and outputs a pair (pc, val_C) or \perp . For any CPU instruction $f_{\text{OP}} \in \mathcal{OP}$, we require that f_{OP} is executable in Bitcoin Script. A *VM program* is an ordered sequence of ℓ elements, denoted $\Pi \in \mathcal{I}^\ell$, where $\mathcal{I} := \{(f_{\text{OP}}, addr_A, addr_B, addr_C) \mid addr_A, addr_B, addr_C \in \mathcal{A}, f_{\text{OP}} \in \mathcal{OP}\}$. We can now define the following.

Definition 6 (VM State). A *VM state*, or simply, *state*, is a triple $S := (M, pc, \Pi)$, where M is the VM memory, pc is the VM program counter, and Π is a VM program.

Definition 7 (State Transition Function). Let $\mathcal{S} := \mathcal{M} \times \mathcal{PC} \times \mathcal{I}^\ell$ be the set of all VM states. We define the *state transition function* $f_{ST} : \mathcal{S} \rightarrow \mathcal{S}$ with f_{ST} taking as argument the state (M_i, pc_i, Π) and giving as output the state (M_{i+1}, pc_{i+1}, Π) as specified in Algorithm 5.

Algorithm 5 State Transition Function f_{ST} .

```

1: function  $f_{ST}(M, pc, \Pi)$ 
2:    $M' \leftarrow M$ ;
3:    $(f_{\text{OP}}, addr_A, addr_B, addr_C) \leftarrow \Pi[pc]$ ;
4:    $val_A \leftarrow M[addr_A]$ ;
5:    $val_B \leftarrow M[addr_B]$ ;
6:    $(pc', val_C) \leftarrow f_{\text{OP}}(val_A, val_B, pc)$ ;
7:   if  $val_C \neq \perp$  then
8:      $M'[addr_C] \leftarrow val_C$ ;
9:   return  $(M', pc', \Pi)$ .

```

⁸ Even though \mathcal{OP} can be arbitrary, we are interested in a Turing-complete instruction set. In particular, we later use ADD, BEQ, and JMP, cf. Algorithm 7 – a well-known Turing-complete instruction set [30].

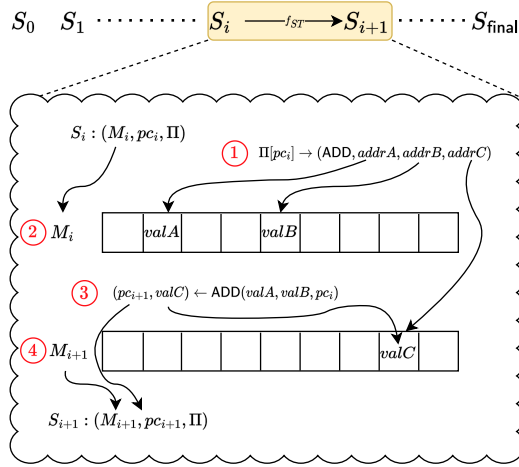


Figure 1. Overview of the state transition function execution f_{ST} . Given a state S_i : (1) instruction $\Pi[pc_i]$ is fetched, (2) values $valA, valB$ are taken from memory at their respective addresses, (3) the instruction is executed, and (4) part of the result (i.e., $valC$) is stored in the memory. The state transition function outputs the new state S_{i+1} .

Given a program Π and a memory configuration M , we assume that the entry point of the program, namely the first instruction that a program executes, is always $\Pi[0]$. Thus, we define as *initial* state the tuple $S_0 := (M, 0, \Pi)$. We use the shorthand notation $f_{ST}^i(S)$ when we apply the state transition function f_{ST} to a state S exactly i times, $f_{ST}^i(S) := f_{ST}(f_{ST}(\dots(f_{ST}(S))))$. We say that a state $S_i := (M_i, pc_i, \Pi)$ at step i is *correct* with respect to an initial state S_0 iff $S_i = f_{ST}^i(S_0)$. We avoid the subscripts (and simply refer to the state S_i as (M, pc, Π)) when it is clear from the context which state we are referring to.

Finally, after a number of execution steps equal to *final* (*final* is decided when a VM instance is created), the program terminates. We denote the final state, or outcome, as $\Pi(S_0) := f_{ST}^{final}(S_0)$. Fig. 1 provides a visual representation of the execution of the state transition function f_{ST} .

We define a VM instance as a tuple

$$\Gamma := \langle \Pi, \text{MemLen}, n, \text{final} \rangle.$$

We write Γ^A to refer to the VM instance executed by party A . We write S_i^A to denote a VM state S_i that A claims to have produced during the execution of A 's VM instance Γ^A . We say that two parties A and B agree on the state S_i if $S_i^A = S_i^B$, and disagree on S_i otherwise.

Definition 8 (Execution Trace Element). Let $(M_i, pc_i, \Pi) := f_{ST}^i(S_0)$, and let MR_i be the root of the Merkle tree with the entries of M_i as its leaves. The i -th VM execution trace element, or simply, i -th trace element is the pair $E_i := (MR_i, pc_i)$, for $i \in \{0, \dots, \text{final}\}$.

We write E_i^A to denote a VM execution trace element E_i that A claims to have produced during the execution of A 's VM instance Γ^A . The VM execution trace is defined as a sequence of consecutive trace elements $ExecTrace :=$

$(E_0, \dots, E_{\text{final}})$. We write $ExecTrace^A$ as a short-hand for $(E_0^A, \dots, E_{\text{final}}^A)$.

We describe how the VM behaves in Algorithm 6: starting from initial state S_0 , it applies the state transition function f_{ST} to the state and records the related trace elements until the program Π ends, namely, once pc is set to be \perp . The VM algorithm is parameterized by *final*, a parameter that represents the maximum number of state transitions that the VM is allowed to perform. The VM algorithm returns as output the VM execution trace $ExecTrace$, along with the resulting memory M after the program execution.

Algorithm 6 The VM algorithm. S_0 is the initial VM state.

```

1: function VMfinal( $S_0$ )
2:    $stepCount \leftarrow 0$ ;
3:   while  $stepCount < \text{final}$  do
4:      $E_{stepCount} \leftarrow (MR, pc)$ ;
5:      $(M, pc, \Pi) \leftarrow f_{ST}(M, pc, \Pi)$ ;
6:     increment  $stepCount$  by 1;
7:    $E_{stepCount} \leftarrow (MR, pc)$ ;
8:    $ExecTrace \leftarrow (E_0, \dots, E_{\text{final}})$ ;
9:   return  $(ExecTrace, M)$ .
```

A practical VM instance. For better readability and to provide a protocol instance that can be deployed in practice, in the rest of the paper, we will consider a VM instance $\Gamma := \langle \Pi, \text{MemLen}, n, \text{final} \rangle$ with the following initialization: We set the length of the memory as $\text{MemLen} = 2^{32}$ and the greatest integer that can be stored in any entry of the memory as $n = 2^{32}$.

Furthermore, we assume that the input program Π has $\ell \leq 2^{32}$ number of instructions⁹ and we set $\text{final} = 2^{32}$.

As for the set \mathcal{OP} of instructions that the VM can execute, our VM instance employs the following: $\mathcal{OP} := \{\text{ADD}, \text{BEQ}, \text{JMP}\}$. This is a minimal set of computer instructions known to be Turing complete [30]. We underscore that the BitVM protocol can function with any Turing-complete instruction set, provided that each instruction within the set is implementable in Bitcoin script. In Algorithm 7, we give an implementation of ADD, BEQ and JMP that can be easily translated in Bitcoin script.

5. The BitVM protocol

The BitVM protocol enhanced the expressiveness of Bitcoin, allowing the encoding of spending conditions based on the outcome of quasi-Turing complete programs.

Protocol overview. The protocol proceeds in four phases.

- 1) In the *setup* phase, P and V agree on the program Π they wish to execute. For example, as described in Section 1, this program could verify the validity of a sequence of chess moves and check whether or not the chess puzzle (encoded in the program itself) has been solved. P and V also agree on the outcome mapping function f , create

9. In the BitVM protocol, we build a Taproot tree where every program instruction is a Tapleaf script. We chose such ℓ since $2^{32} \ll 2^{128}$, the maximum number of leaf scripts in the current specification of Bitcoin [26].

Algorithm 7 The Algorithms ADD, BEQ, and JMP, each taking as input the tuple (pc, val_A, val_B) , and returning a pair (pc, val_C) .

```

1: function ADD( $pc, val_A, val_B$ )
2:   if  $pc = \perp$  then return  $(\perp, \perp)$ ;
3:   return  $(pc + 1, val_A + val_B)$ .

4: function BEQ( $pc, val_A, val_B$ )
5:   if  $pc = \perp$  then return  $(\perp, \perp)$ ;
6:   if  $val_A = val_B$  then
7:     return  $(pc + 1, \perp)$ .
8:   else
9:     return  $(pc + 2, \perp)$ .

10: function JMP( $pc, val_A, val_B$ )
11:   if  $pc = \perp$  then return  $(\perp, \perp)$ ;
12:   return  $(val_A, \perp)$ .

```

and pre-sign all necessary transactions, and post the initial transaction that locks their coins on-chain, which we call Setup.

- 2) In the *VM execute* phase, performed off-chain, P and V generate the input S_0 for Π (e.g., the chess moves starting from a given chessboard state) and compute $\Pi(S_0)$.
- 3) In the *commit* phase, P posts the *CommitComputation* transaction on-chain, i.e., a transaction committing to the input S_0 and the result $\Pi(S_0)$ using Lamport signatures. Upon seeing this, V can either accept the claim as correct and wait for P to settle according to $f(\Pi(S_0))$ (via publishing a *Close* transaction), or, if the committed result does not match $\Pi(S_0)$, initiate a dispute.
- 4) In the event of a dispute, they enter the *resolve dispute* phase. This phase is the main technical challenge of the BitVM protocol, as it requires an on-chain mechanism to verify whether P 's claimed result $\Pi(S_0)$ is correct, while remaining within Bitcoin's scripting limitations.

5.1. Resolving Disputes

The core challenge of BitVM is to enable the on-chain verification of the result of computation that is normally not expressible in Bitcoin Script. To address this, a novel approach is necessary to be able to verify the correct result, $\Pi(S_0)$, when executing Π on input S_0 . Our solution leverages our VM (see Section 4) and the resulting execution trace $ExecTrace := (E_0, \dots, E_{final})$ produced when computing $\Pi(S_0)$. Notably, each successive element in this execution trace results from applying a single VM instruction to the preceding element. We demonstrate that verifying each individual VM instruction can indeed be accomplished within Bitcoin Script, a process we detail later in this section. With this in place, the only remaining task is to identify two consecutive trace elements, where P and V agree on the former but disagree on the latter. Given that they both agree on S_0 , any disagreement over the result $\Pi(S_0)$ ensures that such a pair of consecutive trace elements exists. We identify this pair via an on-chain bisection.

We present both of these components in this section, focusing on describing *what* the parties do by publishing

certain transactions, not *how* exactly these are implemented. The reader can assume that all necessary transaction logic is feasible within Bitcoin, utilizing Lamport signatures and other Bitcoin Script features. We defer the concrete implementation of these Bitcoin transactions and the full protocol specification to Section 6. Whenever we mention that one party must respond with a transaction, this is enforced by a timelock mechanism, allowing the other party to claim all funds in the event of prolonged inactivity by the first party.

In Fig. 2, we illustrate a protocol execution in case of disagreement.

Identify Disagreement. Recall that the total length of the execution trace is final. V initiates the bisection game by publishing a *Kickoff* transaction, forcing P to respond with a *TraceResponse₁* transaction that reveals the middle trace element, $E_{n_{31}}$, where $n_{31} = 2^{31} = \text{final}/2$.

The loop then begins, where V forces P into progressively smaller sequences of the trace by publishing a *TraceChallenge₁* transaction, Lamport committing to a value b_{31} in its witness.

- If V agrees with $E_{n_{31}}$, V commits to $b_{31} = 1$, indicating that P should reveal a trace element in the right half of the sequence.
- If V disagrees with $E_{n_{31}}$, V commits to $b_{31} = 0$, indicating that P should reveal a trace element in the left half of the sequence.
- Based on b_{31} , P responds by publishing a *TraceResponse₂* transaction, revealing the middle element of the next sequence, $E_{n_{30}}$, where $n_{30} = b_{31} \cdot 2^{31} + 2^{30}$.

This process continues recursively, with each step revealing the middle trace element of the current sequence. V publishes a new *TraceChallenge_{32-k}* at each step, setting a new bit b_k (for $k = 30, 29, \dots, 1$), indicating whether to go left ($b_k = 0$) or right ($b_k = 1$). In each response, P commits to the next middle trace element $E_{n_{k-1}}$ in *TraceResponse_{32-k+1}*, for $n_{k-1} = \sum_{i=k}^{31} (b_i \cdot 2^i) + 2^{k-1}$.

With the last transaction, *TraceChallenge₃₂*, V sets the last bit b_0 , indicating (dis)agreement with E_{n_0} , so that we finally obtain the index $\mathcal{N} = \sum_{i=0}^{31} b_i \cdot 2^i$. Let $\mathcal{N}' := \mathcal{N} + 1$. For the pair $(E_{\mathcal{N}}, E_{\mathcal{N}'})$, both P and V agree on the former and disagree on the latter. This pair allows the protocol to resolve the dispute by examining the exact step in the computation where the disagreement occurred.

At this point, P is forced to publish a *CommitInstruction* transaction, committing to the necessary information to execute $S_{\mathcal{N}'} = f_{ST}(S_{\mathcal{N}})$:

- $pc_\theta, pc_{\theta'}$: the program counter of the states $S_{\mathcal{N}}$ and $S_{\mathcal{N}'}$, respectively;
- $insType_\theta \in \mathcal{OP}$: the instruction type at $\Pi[pc_\theta]$;
- $addrA_\theta, addrB_\theta, addrC_\theta$: the memory addresses referenced in $\Pi[pc_\theta]$;
- $valA_\theta, valB_\theta$: the memory values at addresses $addrA_\theta$ and $addrB_\theta$ in $S_{\mathcal{N}}$;
- $valC_\theta$: the value at address $addrC_\theta$ in $S_{\mathcal{N}'}$, i.e., after executing $f_{ST}(S_{\mathcal{N}})$.

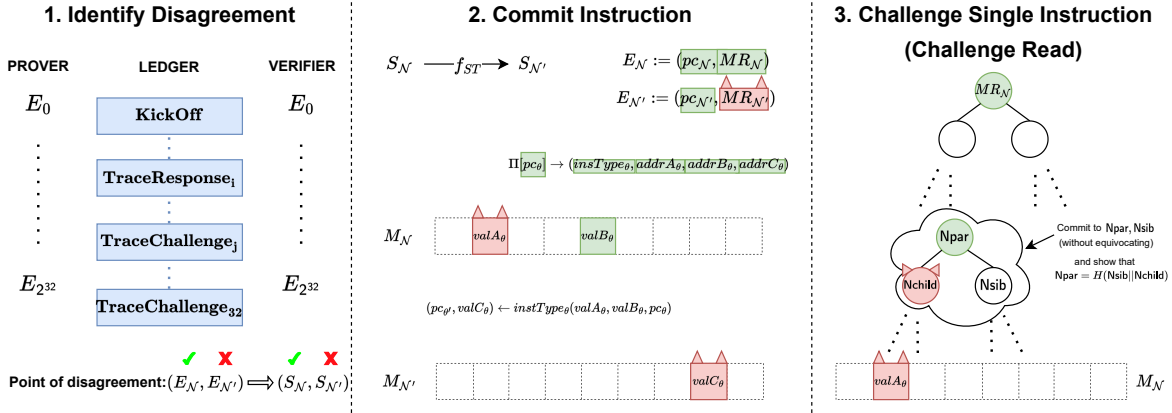


Figure 2. Example of dispute resolution in the BitVM protocol: To resolve a dispute, (1) P and V engage in a bisection game to identify the point of disagreement $(E_N, E_{N'})$ in their execution traces, indicating a disagreement in the transition from state S_N to $S_{N'}$. Next, (2) P commits on-chain to all necessary information for executing $S_{N'} = f_{ST}(S_N)$ (i.e., the values highlighted by colored boxes in the figure). In this example, we assume that P is committing to an incorrect value for $valA_{\theta}$, resulting in incorrect value for $valC_{\theta}$ and $MR_{N'}$. (3) V challenges P through a bisection game over the path in the memory Merkle tree M_N , from the root MR_N to the leaf containing $valA_{\theta}$. This bisection game reveals two intermediate nodes, N_{par} and N_{child} , on which P and V disagree. To get away while using an incorrect value, P would need to prove that N_{child} is indeed the left child of N_{par} , which is impossible without equivocating, leading to punishment.

In addition to committing to these values, the script of $\text{TraceChallenge}_{32}$ requires P to provide them so that $(pc_{\theta'}, valC_{\theta}) = \text{insType}_{\theta}(pc_{\theta}, valA_{\theta}, valB_{\theta})$ holds. In particular, there is a tapleaf that ensures this for each instruction type $\mathcal{OP} := \{\text{ADD}, \text{BEQ}, \text{JMP}\}$.

Challenge Single Instruction. Armed with the prover's commitment to these values and the knowledge that $(pc_{\theta'}, valC_{\theta}) = \text{insType}_{\theta}(pc_{\theta}, valA_{\theta}, valB_{\theta})$ holds, it remains to verify that the committed values correspond to the committed states S_N and $S_{N'}$. There are five potential errors in the committed values: (i) invalid program counter: pc_{θ} is not the program counter of S_N , (ii) invalid program counter next: $pc_{\theta'}$ is not the program counter of $S_{N'}$, (iii) invalid instruction: insType_{θ} , $\text{addr}A_{\theta}$, $\text{addr}B_{\theta}$, or $\text{addr}C_{\theta}$ do not match $\Pi[pc_{\theta}]$, (iv) invalid read: $M_N[\text{addr}A_{\theta}] \neq valA_{\theta}$ or $M_N[\text{addr}B_{\theta}] \neq valB_{\theta}$, and (v) invalid write: $M_{N'}[\text{addr}C_{\theta}] \neq valC_{\theta}$. Referring to Algorithm 5, if $S_{N'} \neq f_{ST}(S_N)$ while $(pc_{\theta'}, valC_{\theta}) = \text{insType}_{\theta}(pc_{\theta}, valA_{\theta}, valB_{\theta})$ holds, one of these five errors must have occurred. In essence, either the values are not correctly derived from the agreed upon state S_N , or the result is not stored correctly in $S_{N'}$. By executing this step, V can pinpoint which of these errors occurred, and thus choose the appropriate challenge path.

Challenging the Current Program Counter. V can challenge the current program counter pc_{θ} by publishing a ChallengeCurrPC transaction. In this transaction, V commits to N , revealing both P 's commitment to pc_N and the currently disputed value pc_{θ} , proving that $pc_N \neq pc_{\theta}$.

Since V already committed to N bit-by-bit during the bisection game with P to locate a disagreement point, V could attempt to equivocate on one of the bits when recommitting to N in the punishment phase. If equivocation occurs, P publishes a PunishCurrPC transaction, revealing

two different bitwise commitments to N linked to the verifier's public key. In this scenario, P claims all funds in the multisignature, thereby terminating the BitVM protocol.

If P fails to act within a specified time interval Δ , and V does not equivocate, V redeems the locked coins, terminating the protocol. This outcome happens either when P remains inactive or when V successfully proves that $pc_N \neq pc_{\theta}$ without equivocating. This case is analogous to challenging the next program counter $pc_{\theta'}$.

Disproving a Wrong Program Instruction. In the CommitInstruction transaction's output locking script, a tapleaf is associated with each line of the program, hard-coding the program line number (pointed to by pc_{θ}) along with the corresponding values insType_{θ} , $\text{addr}A_{\theta}$, $\text{addr}B_{\theta}$, and $\text{addr}C_{\theta}$. If the prover commits to an incorrect program instruction, the verifier can expose this by publishing a DisproveProgram transaction.

To spend one of these tapleaves, V must provide P 's committed values: pc_{θ} , insType_{θ} , $\text{addr}A_{\theta}$, $\text{addr}B_{\theta}$, and $\text{addr}C_{\theta}$. The locking script ensures that the committed program counter pc_{θ} matches the program line number in the tapleaf. If the program counter matches but any other value does not, the DisproveProgram transaction succeeds, and BitVM terminates with V receiving the coins.

Challenge Read. The verifier challenges the correctness of $valA_{\theta}$ at memory address $\text{addr}A_{\theta}$ in S_N , agreeing with the Merkle root MR_N but disputing the value $valA_{\theta}$. To start the challenge, V publishes a ChallengeRead transaction, spending the CommitInstruction output.

The parties engage in a bisection game on the memory Merkle path $\mathcal{P}_R := (MR_N, \dots, M_N[\text{addr}A_{\theta}])$, i.e., the path from the root MR_N to the leaf $M_N[\text{addr}A_{\theta}]$. The goal is to isolate a point of disagreement between V and P . Over 5 rounds of challenge-response (given $\text{MemLen} = 2^{32}$), they progressively narrow down the path until they identify a pair

of nodes ($N_{\text{par}}, N_{\text{child}}$) where they agree on N_{par} (i.e., the parent node) but disagree on N_{child} (i.e., the child node). This bisection game proceeds in a similar fashion as the one for finding the disagreement in the execution trace.

Once this point of disagreement is found, V challenges P to provide a valid sibling node N_{sib} for N_{child} , such that $H(N_{\text{sib}} \parallel N_{\text{child}}) = N_{\text{par}}$. If P cannot provide this due to an invalid commitment, V can claim the locked funds by publishing the `PunishRead` transaction. If P provides the correct sibling and no equivocation occurs, P claims the funds after the timelock expires. This process is analogous for $valB_\theta$ at memory address $addrB_\theta$.

Challenge Write. The process for challenging the result of a write operation is similar to *Challenge Read*, but with a key difference: V challenges the value $valC_\theta$ written to memory at $addrC_\theta$ in $S_{N'}$. This affects both the Merkle path rooted at MR_N and the one rooted at $MR_{N'}$. As a result, the bisection game involves two parallel Merkle paths: \mathcal{P}_W (from MR_N to $M_N[addrC_\theta]$) and \mathcal{P}'_W (from $MR_{N'}$ to $M_{N'}[addrC_\theta]$). In each round, P reveals corresponding nodes from both paths. If P commits to incorrect values on \mathcal{P}_W , V focuses on that path (as in *Challenge Read*). Otherwise, if there is a disagreement on \mathcal{P}'_W , V focuses on that path. The game ends when a pair of parent-child nodes from both paths are isolated, and P must provide a valid sibling node N_{sib} to prove the correctness of the Merkle structure. If P equivocates, V can claim the coins by publishing `PunishWrite`, similarly to *Challenge Read*.

5.2. Honest Closure

In the happy path where the parties agree on the outcome of the computation, the on-chain footprint of the protocol is minimal, with only 3 transactions being published. After a set time period, P can spend the transaction in which they committed to the outcome by signing a pre-signed transaction that includes the result as a witness and distributes the coins according to the outcome mapping function f applied to this result. Should P equivocate, V claims all the funds.

Moreover, BitVM ensures a constant maximum number of transactions in case of dispute, once the VM instance is fixed. For example, consider a VM with 2^{32} memory cells, each capable of storing a 32-bit integer, executing a program up to 2^{32} instructions long, and requiring up to 2^{32} execution steps. Resolving a dispute in the execution of this VM would require publishing at most 81 transactions on-chain. This occurs when V first publishes a `Kickoff` transaction to initiate the `Identify Disagreement` phase. Once this phase is completed, V publishes on-chain a `ChallengeRead` transaction to initiate the `Challenge Read` path (or, similarly, V can initiate the `Challenge Write` path).

6. The BitVM full protocol

In this section, we present the full BitVM protocol specification. All scripts that we use comprise only (multi-)

signature and Lamport signature verification, if/else statements, timelocks, and hashing, and are thus compatible with Bitcoin. Due to space constraints, we present the *setup*, *VM execute*, and *commit* phases and defer the *resolve dispute* phase to Appendix A.

6.1. Setup

In the *setup* phase, the prover P and the verifier V create and presign the necessary transactions for both honest protocol execution and potential dispute resolution; then both P and V lock an on-chain deposit, in_P and in_V , respectively.

At first, both P and V create all the transactions that are defined in this section and Appendix A, except `Setup`. Whenever such a transaction contains a new Lamport public key, the corresponding party creates one using `Lamp.KeyGen` and shares the public key with the other party.

Each transaction output either requires a 2-of-2 multisignature σ_{PV} to be spent and is presigned by both parties or requires a signature from one party along with a timelock. The timelock condition ensures that if a party ceases participation in the BitVM protocol, they forfeit the deposit, which the counterparty can then claim, along with their deposit.

After creating the transactions, the parties exchange them for presigning. For each transaction P (V) verifies it is well-formed according to the definitions below. If verified, the transaction is signed and sent to V (P). The

Finally, P and V sign and publish the `Setup` (cf. Eq. (1)) transaction on-chain. This transaction takes inputs from P (holding in_P) and V (holding in_V), creating an output that can be unlocked by both parties under the following conditions: P can spend the output by presenting Lamport commitments to MR_0 and MR_{final} (i.e., the commitments to the input and output of program Π), while V can spend the output after a timeout period Δ . For brevity, we use σ_U when the signature is in a transaction's witness; the message signed in that case is the transaction body (inputs and outputs without witnesses).

$\text{Setup} :=$

$$\begin{aligned} (in &= [(*, *, \text{CheckSig}_{pk_P}), (*, *, \text{CheckSig}_{pk_V})], \\ wit &= [(\sigma_P), (\sigma_V)], \\ out &= [(d\mathbb{B}; \langle \text{CommitComputationScript}, \text{TL}(\Delta) \wedge \\ &\quad \text{CheckSig}_{pk_V} \rangle)]), \end{aligned} \quad (1)$$

The script `CommitComputationScript` is defined below.

$\text{CommitComputationScript} :=$

$$\text{CheckMSig}_{pk_{PV}} \wedge \text{CheckComm}_{pk_{E_0}} \wedge \text{CheckComm}_{pk_{E_{\text{final}}}}.$$

6.2. VM Execute

The prover P sends to the verifier V the input x of program Π via a communication channel. Both P and V execute off-chain the program Π with input x on their

VM instance. They copy x into the VM memory M and call Algorithm 6 with input $S_0 := (M, 0, \Pi)$. They get as output the VM execution trace $ExecTrace$ and the memory M , from which they fetch the output y of program Π with input x . We stress that this is the most resource-intensive phase of BitVM and it is entirely performed *off-chain*.

6.3. Commit

The prover P publishes the `CommitComputation` transaction (cf. Eq. (2)) on-chain, which spends the output of the `Setup` transaction by providing a Lamport commitment to $E_0 := (MR_0, pc_0)$ and $E_{final} := (MR_{final}, pc_{final})$.

$$\begin{aligned} \text{CommitComputation} := & \\ & ([in = (\text{Setup}, 0, \text{CommitComputationScript})], \\ & [wit = (\sigma_{PV}, E_0, c_{E_0}, E_{final}, c_{E_{final}})], \\ & [out = (d\$, \langle \text{CheckMSig}_{\text{pk}_{PV}}, \text{CloseScript}, \\ & \quad \text{CheckSig}_{\text{pk}_V} \wedge \text{TL}(2\Delta) \rangle)]). \end{aligned} \quad (2)$$

The script `CloseScripti` is defined in Algorithm 8.

Algorithm 8 The script `CloseScripti`. In the setup phase, the public key $\text{pk}_{MR_{final}}$ is hard-coded in the script.

```

1: function CloseScripti( $\sigma_{PV}$ ,  $MR_{final}$ ,  $c_{MR_{final}}$ )
2:    $\text{TL}(\Delta)$ ;
3:    $\text{CheckMSigVerify}_{\text{pk}_{PV}}(\sigma_{PV})$ ;
4:    $\text{CheckCommVerify}_{\text{pk}_{MR_{final}}} (MR_{final}, c_{MR_{final}})$ ;
5:   if  $MR_{final} = \mathcal{MR}_i$  then
6:     return True;
7:   return False.

```

The verifier V can either challenge P if they disagree with the MR_{final} published on-chain by P or simply take no action if they agree. Since the VM execution is deterministic, honest parties running the same program on the same input naturally agree on MR_{final} . A disagreement, therefore, implies that one party is behaving dishonestly.

Close. V agrees with P 's commitment to MR_{final} and does not dispute it. The BitVM protocol follows the happy path: after a timeout period Δ , P publishes one of the close transactions, $\text{Close}_1, \dots, \text{Close}_m$. Each of these transactions distributes the funds according to the outcome mapping function f , applied to one of the possible results of the computation¹⁰. If P does not publish any Close_i transaction after that $\text{TL}(2\Delta)$ expires after the publication of `CommitComputation` transaction, V can unlock `CommitComputation` output with their signature and claim all the funds.

Transaction Close_i (cf. Eq. (3)) spends the output of `CommitComputation` by unlocking `CloseScripti` and creates two outputs. The first output carries $o_P\$$ and can be

10. During the setup phase, P and V agree on f and jointly create and sign a finite set of closing transactions, one for each possible outcome. The funds are distributed to P and V according to the result of f .

unlocked by P after a timeout period Δ or by V if P equivocates on MR_{final} (as shown in Algorithm 9). The second output carries $o_V\$$ and can be unlocked by V .

$$\begin{aligned} \text{Close}_i := & \\ & ([in = (\text{CommitComputation}, 0, \text{CloseScript})], \\ & [wit = (\sigma_{PV}, MR_{final}, c_{MR_{final}})], \\ & [out = (v_P; \langle \text{CheckSig}_{\text{pk}_P} \wedge \text{TL}(\Delta), \\ & \quad \text{PunishCloseScript} \rangle), (v_V; \text{CheckSig}_{\text{pk}_V})]). \end{aligned} \quad (3)$$

Algorithm 9 The script `PunishCloseScript`. In the setup phase, the public key $\text{pk}_{MR_{final}}$ is hard-coded in the script.

```

1: function PunishCloseScript( $\sigma_{PV}$ ,  $c_0$ ,  $c_1$ )
2:    $\text{CheckMSigVerify}_{\text{pk}_{PV}}(\sigma_{PV})$ ;
3:   for  $i = 1, \dots, |MR_{final}|_{\text{bit}}$  do
4:     if  $\text{Equivocation}(\text{pk}_{MR_{final}[i]}, c_0, c_1) = \text{True}$  then
5:       return True;
6:   return False.

```

Identify Disagreement. V disagrees with P 's commitment to MR_{final} . To dispute P 's result, V publishes the `KickOff` transaction (cf. Eq. (4)) by spending `CommitComputation`'s output, unlocking it through the multisignature.

$$\begin{aligned} \text{KickOff} := & \\ & ([in = (\text{CommitComputation}, 0, \text{CheckMSig}_{\text{pk}_{PV}})], \\ & [wit = [(\sigma_{PV})], \\ & [out = [(d\$, \langle \text{ChallScript}_1, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_V} \rangle)]]). \end{aligned} \quad (4)$$

The script `ChallScriptj`, with $j \in \{1, \dots, 31\}$, is defined as follows:

$$\text{ChallScript}_j := \text{CheckMSig}_{\text{pk}_{PV}} \wedge \text{CheckComm}_{\text{pk}_{E_{n_{32-j}}}}.$$

The parties engage in an on-chain interactive protocol known as dispute bisection game (cf. [31, Appendix B.1]): the game is played over the VM execution trace $ExecTrace := (E_0, \dots, E_{final})$ and has the goal to determine a pair of consecutive VM trace elements $(E_{\mathcal{N}}, E_{\mathcal{N}'})$, where $\mathcal{N}' := \mathcal{N} + 1$, such that they agree on $E_{\mathcal{N}}$ and disagree on $E_{\mathcal{N}'}$.

After that, V initiates the bisection game by publishing the `Kickoff` transaction, P responds by publishing the `TraceResponse1` transaction (cf. Eq. (5)), committing to $E_{n_{31}}$ in the witness, where $n_{31} = 1 \cdot 2^{31}$.

$$\begin{aligned} \text{TraceResponse}_1 := & \\ & ([in = [(\text{KickOff}, 0, \text{CheckMSig}_{\text{pk}_{PV}} \wedge \\ & \quad \text{CheckComm}_{\text{pk}_{E_{n_{31}}}})], \\ & [wit = [(\sigma_{PV}, E_{n_{31}}, c_{E_{n_{31}}})], \\ & [out = [(d\$, \langle \text{RespScript}_1, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_P} \rangle)]]). \end{aligned} \quad (5)$$

The script RespScript_i , with $i \in \{1, \dots, 32\}$, is defined as follows:

$$\text{RespScript}_i := \text{CheckMSig}_{\text{pk}_{PV}} \wedge \text{CheckComm}_{\text{pk}_{b_{32-i}}}.$$

Next, V publishes the TraceChallenge_1 transaction (cf. Eq. (6)), committing to bit b_{31} in the witness.

$$\begin{aligned} \text{TraceChallenge}_1 := & \\ (in = [(\text{TraceResponse}_1, 0, \text{RespScript}_1)], & \\ wit = [(\sigma_{PV}, b_{31}, c_{b_{31}})], & \\ out = [(d\mathbb{B}; \langle \text{ChallScript}_2, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_V} \rangle)]) & \end{aligned} \quad (6)$$

During the dispute bisection game, P publishes transactions TraceResponse_i (cf. Eq. (7)), with $i = 1, \dots, 32$, and V publishes transactions TraceChallenge_j (cf. Eq. (8)), with $j = 1, \dots, 31$.

$$\begin{aligned} \text{TraceResponse}_i := & \\ (in = [(\text{TraceChallenge}_{i-1}, 0, \text{ChallScript}_{i-1})], & \\ wit = [(\sigma_{PV}, E_{n_{32-i}}, c_{E_{n_{32-i}}})], & \\ out = [(d\mathbb{B}; \langle \text{RespScript}_i, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_P} \rangle)]) & \end{aligned} \quad (7)$$

$$\text{where } n_{32-i} = 1 \cdot 2^{32-i} + \sum_{k=32-i+1}^{31} b_k \cdot 2^k.$$

$$\begin{aligned} \text{TraceChallenge}_j := & \\ (in = [(\text{TraceResponse}_j, 0, \text{RespScript}_j)], & \\ wit = [(\sigma_{PV}, b_{32-j}, c_{b_{32-j}})], & \\ out = [(d\mathbb{B}; \langle \text{ChallScript}_{j+1}, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_V} \rangle)]) & \end{aligned} \quad (8)$$

Finally, V publishes $\text{TraceChallenge}_{32}$ (cf. Eq. (9)).

$$\begin{aligned} \text{TraceChallenge}_{32} := & \\ (in = [(\text{TraceResponse}_{32}, 0, \text{RespScript}_{32})], & \\ wit = [(\sigma_{PV}, b_0, c_{b_0})], & \\ out = [(d\mathbb{B}; \langle \text{ADDScript}, \text{BEQScript}, \text{JMPScript}, & \\ \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_V} \rangle)]) & \end{aligned} \quad (9)$$

To unlock the $\text{TraceChallenge}_{32}$ output, P is forced to provide a commitment for $pc_\theta, pc_{\theta'}, insType_\theta, addrA_\theta, addrB_\theta, addrC_\theta, valA_\theta, valB_\theta, valC_\theta$. The instruction $insType_\theta$ must match the leaf being spent and $pc_\theta, pc_{\theta'}, valA_\theta, valB_\theta, valC_\theta$ must align with the instruction's semantics. For instance, if P unlocks the ADDScript tapleaf (cf. Algorithm 10), the condition $\text{ADD}(pc_\theta, valA_\theta, valB_\theta) = (pc_{\theta'}, valC_\theta)$ must hold, where ADD is the VM instruction defined in Algorithm 7, lines 1 to 3. The leaves BEQScript and JMPScript are analogous to ADDScript but they encode the semantics of the BEQ and JMP instructions, respectively. The *resolve dispute* phase is deferred to Appendix A.

Algorithm 10 The script ADDScript . In the setup phase, the public keys $\text{pk}_{pc_\theta}, \text{pk}_{pc_{\theta'}}, \text{pk}_{insType_\theta}, \text{pk}_{addrA_\theta}, \text{pk}_{addrB_\theta}, \text{pk}_{addrC_\theta}, \text{pk}_{valA_\theta}, \text{pk}_{valB_\theta}, \text{pk}_{valC_\theta}$ and the semantics of the ADD instruction are hard-coded in the script.

```

1: function  $\text{ADDScript}(\sigma_{PV}, pc_\theta, c_{pc_\theta}, pc_{\theta'}, c_{pc_{\theta'}}, insType_\theta,$ 
    $c_{insType_\theta}, addrA_\theta, c_{addrA_\theta}, addrB_\theta, c_{addrB_\theta}, addrC_\theta,$ 
    $c_{addrC_\theta}, valA_\theta, c_{valA_\theta}, valB_\theta, c_{valB_\theta}, valC_\theta, c_{valC_\theta})$ 
2:    $\text{CheckMSigVerify}_{\text{pk}_{PV}}(\sigma_{PV});$ 
3:    $\text{CheckCommVerify}_{\text{pk}_{pc_\theta}}(pc_\theta, c_{pc_\theta});$ 
4:    $\text{CheckCommVerify}_{\text{pk}_{pc_{\theta'}}}(pc_{\theta'}, c_{pc_{\theta'}});$ 
5:    $\text{CheckCommVerify}_{\text{pk}_{insType_\theta}}(insType_\theta, c_{insType_\theta});$ 
6:    $\text{CheckCommVerify}_{\text{pk}_{addrA_\theta}}(addrA_\theta, c_{addrA_\theta});$ 
7:    $\text{CheckCommVerify}_{\text{pk}_{addrB_\theta}}(addrB_\theta, c_{addrB_\theta});$ 
8:    $\text{CheckCommVerify}_{\text{pk}_{addrC_\theta}}(addrC_\theta, c_{addrC_\theta});$ 
9:    $\text{CheckCommVerify}_{\text{pk}_{valA_\theta}}(valA_\theta, c_{valA_\theta});$ 
10:   $\text{CheckCommVerify}_{\text{pk}_{valB_\theta}}(valB_\theta, c_{valB_\theta});$ 
11:   $\text{CheckCommVerify}_{\text{pk}_{valC_\theta}}(valC_\theta, c_{valC_\theta});$ 
12:  if  $insType_\theta = \text{ADD} \wedge \text{ADD}(pc_\theta, valA_\theta, valB_\theta) =$ 
    $(pc_{\theta'}, valC_\theta)$  then
13:    return True;
14:  else
15:    return False.

```

7. Security Analysis

In this section, we show that BitVM is an *on-chain state verification protocol* that satisfies *balance security* and *rational correctness* under our model. Due to space constraints, we only provide proof sketches and defer the full proofs to the extended version of the paper [31], where we model BitVM as an Extensive Form Game (EFG).

Theorem 7.1. BitVM is an on-chain state verification protocol that achieves *balance security* and *rational correctness* as defined in Definition 5.

7.1. Balance Security

We consider two scenarios: (i) both parties act honestly, and (ii) one party, $A \in \{P, V\}$, deviates at any step of the protocol. In both cases, we prove that an honest party retains their funds.

Setup. If either party deviates during Setup, the honest party will refuse to sign the Setup transaction, ensuring no coins are locked unless both parties have received all necessary pre-signed transactions (cf. [31, Lemma D.1]).

Both parties honest. When both parties are honest, BitVM follows an optimistic path: P posts the correct computation result on-chain in CommitComputation and, after a timelock Δ , publishes Close to distribute funds according to the outcome function f (cf. [31, Lemma D.3]).

V is honest and P is Byzantine. If P fails to publish CommitComputation , because of either being inactive or executing invalid computations, or subsequently fails to post Close , V can claim the coins after the relevant timelocks

expire (Δ or 2Δ) (cf. [31, Lemmas D.2, D.3]). This mechanism prevents hostage scenarios by enabling V to reclaim funds in case of non-responsiveness.

If P has committed an incorrect computation result in `CommitComputation`, V publishes `KickOff` (Identify Disagreement phase). As shown in (cf. [31, Lemma D.4]), if P is inactive, V can claim the coins after the relevant timelock expires. If the Identify Disagreement phase completes, V knows a step for which P has committed on-chain to an execution of the VM state transition function (Algorithm 5) incorrectly (cf. [31, Lemma D.7]).

P can deviate in the following ways: by using an incorrect program counter (current or next), making an incorrect memory read or write, or using invalid instructions. For each of these deviations, V can post the corresponding transaction on-chain (e.g., `ChallengeCurrPC`, `ChallengeNextPC`, `ChallengeRead`, `ChallengeWrite`, `PunishInstruction`). Following the respective dispute path, V is able to disprove P 's computation and claim the coins, as we conclude in [31, Lemma D.14]).

P is honest and V is Byzantine. V can misbehave by publishing `KickOff` on-chain to initiate the Dispute Phase, although P has committed to the correct output of the execution in `CommitComputation`. We show that if V is inactive during the Identify Disagreement phase, P can claim the coins after the timelock expires (cf. [31, Lemma D.5]). Otherwise, if the Identify Disagreement phase completes, V has to publish one of the transactions `ChallengeCurrPC`, `ChallengeNextPC`, `ChallengeRead`, `ChallengeWrite`, `PunishInstruction`. Since P has committed to the correct values, V cannot disprove P , as we conclude in [31, Lemma D.15]).

7.2. Rational Correctness

We have shown that if a party $A \in \{P, V\}$ misbehaves at any point, the other party A' can claim all the coins. Therefore, when both parties are rational, no party will misbehave but instead follow the optimistic path of BitVM.

8. Discussion

In the following, we outline BitVM's costs to demonstrate its feasibility and efficiency, discuss potential applications, its limitations – particularly its permissioned nature – and conclude with possible improvements.

Feasibility & Cost Evaluation. To assess the feasibility of our approach, we estimate the transaction fees for both an optimistic run and the most expensive dispute run of BitVM, using the VM instance defined in Section 4.

We assume constant transaction fees of $3sat/vB^{11}$, a Bitcoin price of 64,300\$ (as of 24 September 2024). For the Lamport signatures, we set $\lambda = 160$, thus public key

and commitment are of length $20B$; for an a -bit message they occupy $a \times 20B = a \times 5vB$. We also assume that $\Delta = 12$ hours, meaning each timelock expires after half a day. Different concrete values can be chosen, but any such selection would require scaling the evaluation accordingly. For details on transaction size calculations, we refer the reader to the extended version [31, Appendix E].

Optimistic case. In the optimistic case, three on-chain transactions are published: `Setup` ($187vB$), `CommitComputation` ($2,093vB$), and `Close` ($1,289vB$), totaling $3,569vB$. The protocol's execution cost is 10,707 sat (6.90\$). In terms of execution time, once `Setup` is published, BitVM completes in at most $2u$ rounds.

Dispute case. We focus on the most expensive path in terms of fees, the *Challenge Write* path. Besides the `Setup` and the `CommitComputation` transactions, the following transactions are also published on-chain in case of dispute:

- 1 `KickOff`, 32 `TraceChallenge`, 32 `TraceResponse`, 1 `ChallengeWrite`, 5 `WriteChallenge` transactions, with up to $1,112vB$.
- 5 `WriteResponse`, transactions, with up to $2,093vB$.
- The transactions `CommitWrite1` and `PunishWrite1` with up to $7,286vB$.
- The transaction `CommitInstruction` with up to $2,751vB$.

Overall, the path weighs $109,020vB$. The total protocol execution cost is roughly $327Ksat$, or about 210\$, and, once the `Setup` transaction is published on-chain, it takes at most $80 \times \Delta = 40$ days to complete its execution.

Applications, Limitations & Improvements. This paper aims to formalize BitVM and prove its security, establishing for the first time that quasi-Turing complete on-chain state verification is feasible on blockchains like Bitcoin. By facilitating the verification of such programs on Bitcoin, BitVM unlocks a diverse array of applications, including cross-chain bridges, light clients, zk-proof verification, state channels, games, and escrow contracts. Off-chain computation by users eliminates additional costs for miners, making BitVM a practical and scalable solution for complex applications.

Nevertheless, the current construction is not optimized for efficiency. Possible improvements include undergoing engineering efforts [15], [16], replacing Lamport signatures with Winternitz signatures [28], or incorporating future op-codes, such as covenants [2].

Furthermore, the current construction is a two-party permissioned protocol, limiting its direct applicability compared to [19], for example, which permits anyone to act as the verifier. However, this added flexibility incurs a significantly higher cost of $1.9K\$$ (due to filling an entire 4MB Bitcoin block), compared to BitVM's practical cost range of $7 - 200\$$.

11. In Bitcoin, the size of a SegWit [25] transaction is expressed in *virtual Bytes*, or *vBytes* (vB). The number of vBytes of a transaction witness is equal to its number of Bytes divided by four.

References

- [1] “Bitcoin script,” <https://en.bitcoin.it/wiki/Script>, accessed: 2024-11.
- [2] M. Bartoletti, S. Lande, and R. Zunino, “Bitcoin covenants unchained,” in *Leveraging Applications of Formal Methods, Verification and Validation: Applications*, 2020, pp. 25–42.
- [3] M. Bartoletti, R. Marchesin, and R. Zunino, “Secure compilation of rich smart contracts on poor UTXO blockchains,” in *9th IEEE European Symposium on Security and Privacy, EuroS&P*, 2024. [Online]. Available: <https://doi.org/10.1109/EuroSP60621.2024.00021>
- [4] “Mate in 1: white to move,” https://lichess.org/analysis/1Bb3BN/R2Pk2r/1Q5B/4q2R/2bN4/4Q1BK/1p6/1bq1R1rb_w_-.
- [5] “Mate in 1: white to move,” https://lichess.org/analysis/r1qr1b2/1R3pkp/P2p2pN/ppnPp1Q1/bn2P3/4P1NP/PBBPp1P1/5RK1_w_Q_e6_0_1.
- [6] M. Bartoletti and R. Zunino, “Bitml: A calculus for bitcoin smart contracts,” in *ACM CCS*, 2018, p. 83–100. [Online]. Available: <https://doi.org/10.1145/3243734.3243795>
- [7] M. Bartoletti, T. Cimoli, and R. Zunino, “Fun with bitcoin smart contracts,” in *Leveraging Applications of Formal Methods, Verification and Validation. Industrial Practice*, Cham, 2018, pp. 432–449.
- [8] P. Das, L. Eckey, T. Frassetto, D. Gens, K. Hostáková, P. Jauernig, S. Faust, and A.-R. Sadeghi, “FastKitten: Practical smart contracts on bitcoin,” in *USENIX Security 19*, 2019.
- [9] T. Frassetto, P. Jauernig, D. Koisser, D. Kretzler, B. Schlosser, S. Faust, and A.-R. Sadeghi, “Pose: Practical off-chain smart contract execution,” in *NDSS Symposium*. Internet Society, 2023. [Online]. Available: <http://dx.doi.org/10.14722/ndss.2023.23118>
- [10] Ethereum Foundation, “State channels,” 2024, accessed: 2024-11. [Online]. Available: <https://ethereum.org/en/developers/docs/scaling/state-channels/>
- [11] S. Dziembowski, S. Faust, and K. Hostáková, “General state channel networks,” in *ACM CCS*, 2018, p. 949–966. [Online]. Available: <https://doi.org/10.1145/3243734.3243856>
- [12] L. Aumayr, O. Ersoy, A. Erwig, S. Faust, K. Hostáková, M. Maffei, P. Moreno-Sanchez, and S. Riahi, “Generalized channels from limited blockchain scripts and adaptor signatures,” in *ASIACRYPT*, 2021.
- [13] C. DLC, “Discreet log contracts: An overview,” 2018. [Online]. Available: <https://adiabat.github.io/dlc.pdf>
- [14] V. Madathil, S. A. Thyagarajan, D. Vasilopoulos, L. Fournier, G. Malavolta, and P. Moreno-Sanchez, “Cryptographic oracle-based conditional payments,” in *NDSS Symposium*, 2023. [Online]. Available: <https://dx.doi.org/10.14722/ndss.2023.24024>
- [15] FairGate Labs, “Bitvmx & bitvm,” 2024, accessed: 2024-11. [Online]. Available: <https://fairgate.io>
- [16] Sovryn, “Bitcoins,” 2024, accessed: 2024-11. [Online]. Available: <https://sovryn.com/bitcoins>
- [17] Citrea, “Bitcoin settlement trust-minimized btc bridge: Bitvm,” 2024, accessed: 2024-11. [Online]. Available: <https://docs.citrea.xyz/technical-specs/characteristics/bitcoin-settlement-trust-minimized-btc-bridge/bitvm>
- [18] Robin Linus, “Bitvm: Compute anything on bitcoin,” 2023, accessed: 2024-11. [Online]. Available: <https://bitvm.org/bitvm.pdf>
- [19] R. Linus, L. Aumayr, A. Zamyatin, A. Pelosi, Z. Avarikioti, and M. Maffei, “BitVM2: Bridging bitcoin to second layers,” 2024, accessed: 2024-11. [Online]. Available: https://bitvm.org/bitvm_bridge.pdf
- [20] J. Garay, A. Kiayias, and N. Leonardos, “The bitcoin backbone protocol: Analysis and applications,” *J. ACM*, vol. 71, no. 4, Aug. 2024. [Online]. Available: <https://doi.org/10.1145/3653445>
- [21] A. Tzinas, S. Sridhar, and D. Zindros, “On-chain timestamps are accurate,” *Cryptology ePrint Archive*, Paper 2023/1648, 2023. [Online]. Available: <https://eprint.iacr.org/2023/1648>
- [22] L. Aumayr, Z. Avarikioti, M. Maffei, G. Scaffino, and D. Zindros, “Blink: An optimal proof of proof-of-work,” *Cryptology ePrint Archive*, 2024. [Online]. Available: <https://eprint.iacr.org/2024/692>
- [23] J. Poon and T. Dryja, “The bitcoin lightning network: Scalable off-chain instant payments,” 2016.
- [24] S. Rain, G. Avarikioti, L. Kovács, and M. Maffei, “Towards a game-theoretic security analysis of off-chain protocols,” in *IEEE CSF*, 2023, pp. 107–122. [Online]. Available: <https://doi.org/10.1109/CSF57540.2023.00003>
- [25] E. Lombrozo and P. Wuille, “Bip 141, segregated witness,” 2015, accessed: 2024-11. [Online]. Available: <https://github.com/bitcoin/bips/blob/master/bip-0141.mediawiki>
- [26] P. Wuille, J. Nick, and A. Towns, “Bip 0341, taproot: Segwit version 1 spending rules,” jan 2020, https://en.bitcoin.it/wiki/BIP_0341.
- [27] L. Lamport, “Constructing digital signatures from a one way function,” *Tech. Rep. CSL-98*, October 1979, <https://www.microsoft.com/en-us/research/publication/constructing-digital-signatures-one-way-function/>.
- [28] D. Boneh and V. Shoup, “A graduate course in applied cryptography,” *Draft 0.6*, 2023, <https://toc.cryptobook.us/>.
- [29] ZeroSync, “BitVM Github repository,” dec 2023, <https://github.com/BitVM/BitVM>.
- [30] Esolangs, “Addleq,” <https://esolangs.org/wiki/Addleq>.
- [31] “Bitvm: Quasi-turing complete computation on bitcoin (extended version),” <https://bitvm-paper.github.io/>.

Appendix A.

The BitVM Full Protocol: Resolve Dispute

A.1. Resolve Dispute

P spends the $\text{TraceChallenge}_{32}$ output by publishing the CommitInstruction transaction (cf. Eq. (10)).

$$\begin{aligned}
 \text{CommitInstruction} := & \\
 & \left(in = [(\text{TraceChallenge}_{32}, 0, \text{OPScript})], \right. \\
 & \quad wit = [(\sigma_{PV}, pc_\theta, c_{pc_\theta}, pc_{\theta'}, c_{pc_{\theta'}}, insType_\theta, \\
 & \quad \quad c_{insType_\theta}, addrA_\theta, c_{addrA_\theta}, addrB_\theta, c_{addrB_\theta}, addrC_\theta, \\
 & \quad \quad c_{addrC_\theta}, valA_\theta, c_{valA_\theta}, valB_\theta, c_{valB_\theta}, valC_\theta, c_{valC_\theta}), \\
 & \quad out = [(d\mathbb{B}; \langle \text{CheckMSig}_{pk_{PV}}, \{\text{CIScriptPCCurr}_i\}_{i \in \{1, \dots, 32\}}, \\
 & \quad \quad \{\text{CIScriptPCNext}_i\}_{i \in \{1, \dots, 32\}}, \{\text{CIScriptInstr}_j\}_{j \in \{1, \dots, \ell\}}, \\
 & \quad \quad \text{TL}(\Delta) \wedge \text{CheckSig}_{pk_P})] \Big). \tag{10}
 \end{aligned}$$

The tableaf that P unlocks when publishing CommitInstruction is $\text{OPScript} \in \{\text{ADDScript}, \text{BEQScript}, \text{JMPScript}\}$.

By publishing the CommitInstruction transaction, P reveals all the information necessary for the state transition from S_N to $S_{N'}$. Depending on the specific error that V claims P made, V spends the output of CommitInstruction in one of the following ways.

A.1.1. Challenging the Current Program Counter. V is claiming that, by publishing CommitInstruction , P is committing to a program counter pc_θ at step N that differs from the program counter pc_N (previously committed by P during the dispute bisection game). V challenges the

current program counter pc_θ by unlocking one of the leaves CIScriptPCCurr_i (cf. Algorithm 11) via the publication of transaction ChallengeCurrPC (cf. Eq. (11)). We use Algorithm 12 to map the challenge-response rounds to the leaves $\text{CIScriptPCCurr}_0, \dots, \text{CIScriptPCCurr}_{31}$. When V unlocks leaf CIScriptPCCurr_i , they challenge the program counter of the $(32 - i)$ -th challenge-response round of the dispute bisection game.

Algorithm 11 The script CIScriptPCCurr_i , for $i \in \{0, \dots, 31\}$. For each CIScriptPCCurr_i , in the setup phase, we hard-code the public keys $\text{pk}_{pc_\theta}, \text{pk}_{\mathcal{N}}$. For each CIScriptPCCurr_i , for $i \in \{1, \dots, 31\}$, we hard-code the same public key pk_{pc_i} hard-coded in ChallScript_i . For CIScriptPCCurr_0 , we hard-code the same public key pk_{pc_0} hard-coded in $\text{CommitComputationScript}$.

```

1: function  $\text{CIScriptPCCurr}_i(\sigma_{PV}, \mathcal{N}, c_{\mathcal{N}}, pc_i, c_{pc_i}, pc_\theta, c_{pc_\theta})$ 
2:    $\text{CheckMSigVerify}_{\text{pk}_{PV}}(\sigma_{PV});$ 
3:    $\text{CheckCommVerify}_{\text{pk}_{\mathcal{N}}}(\mathcal{N}, c_{\mathcal{N}});$ 
4:   if  $\text{CountZeroes}(\mathcal{N}) \neq i$  then
5:      $\triangleright$  Maps  $\mathcal{N}$  to one of the 32 program counters  $pc_{n_0}, \dots, pc_{n_{31}}$ .  $\triangleleft$ 
6:     return False;
7:    $\text{CheckCommVerify}_{\text{pk}_{pc_i}}(pc_i, c_{pc_i});$ 
8:    $\text{CheckCommVerify}_{\text{pk}_{pc_\theta}}(pc_\theta, c_{pc_\theta});$ 
9:   if  $pc_i \neq pc_\theta$  then
10:    return True;
11:   else
12:    return False.

```

Algorithm 12 The algorithm CountZeroes . It counts the number of consecutive bits set to 0 in the binary representation of a number N , starting from the least significant bit (LSB), until the first occurrence of a bit set to 1.

```

1: function  $\text{CountZeroes}(N)$ 
2:    $\text{counter} \leftarrow 0;$ 
3:    $\text{flag} \leftarrow \text{False};$ 
4:   for  $i = 0, \dots, |N|_{\text{bit}} - 1$  do
5:     if  $N[|N|_{\text{bit}} - i] = 1$  then
6:        $\text{flag} \leftarrow \text{True};$ 
7:        $\triangleright$  Set the flag, stop incrementing the counter.  $\triangleleft$ 
8:     else
9:       if  $\text{flag} = \text{False}$  then
10:         $\text{counter} \leftarrow \text{counter} + 1;$ 
11:   return counter.

```

$\text{ChallengeCurrPC} :=$

$$\left(in = [(\text{CommitInstruction}, 0, \text{CIScriptPCCurr}_{\mathcal{N}})], \right.$$

$$wit = [(\sigma_{PV}, \mathcal{N}, c_{\mathcal{N}}, pc_{\mathcal{N}}, c_{pc_{\mathcal{N}}}, pc_\theta, c_{pc_\theta})],$$

$$out = [(d\$, \langle \text{ChallPCScript}, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_V} \rangle)] \Big). \quad (11)$$

In the ChallengeCurrPC transaction, V commits again to \mathcal{N} , potentially equivocating. P can punish equivocation by unlocking ChallPCScript script (cf. Algorithm 13).

If V equivocates, P publishes PunishCurrPC (cf. Eq. (12)), redeeming all the funds in the multisignature.

Algorithm 13 The script ChallPCScript . In the setup phase, the public key $\text{pk}_{\mathcal{N}}$ is hard-coded in the script.

```

1: function  $\text{ChallPCScript}(\sigma_{PV}, c_0, c_1)$ 
2:    $\text{CheckMSigVerify}_{\text{pk}_{PV}}(\sigma_{PV});$ 
3:   for  $i = 1, \dots, |\mathcal{N}|_{\text{bit}}$  do
4:     if  $\text{Equivocation}(\text{pk}_{\mathcal{N}[i]}, c_0, c_1) = \text{True}$  then
5:       return True;
6:   return False.

```

$\text{PunishCurrPC} :=$

$$\left(in = [(\text{ChallengeCurrPC}, 0, \text{ChallPCScript})], \right.$$

$$wit = [(\sigma_{PV}, c_0, c_1)],$$

$$out = [(d\$, \text{CheckSig}_{\text{pk}_P})] \Big). \quad (12)$$

Challenging the next program counter works analogous, thus deferred to the extended version of this paper [31].

A.1.2. Punish Wrong Instruction. P has committed to a current program counter pc_θ that does not correspond to the correct program instruction, specifically: $\Pi[pc_\theta] \neq (\text{insType}_\theta, \text{addrA}_\theta, \text{addrB}_\theta, \text{addrC}_\theta)$.

V spends the CommitInstruction output by unlocking the script CIScriptInstr_j (cf. Algorithm 14) and publishing the DisproveProgram transaction (cf. Eq. (13)). A script CIScriptInstr exists for each of the ℓ instructions in the program Π .

Algorithm 14 The script CIScriptInstr_j , for $j \in \{1, \dots, \ell\}$. In the script CIScriptInstr_j , during the setup phase we hard-code the public keys $\text{pk}_{pc_\theta}, \text{pk}_{\text{insType}_\theta}, \text{pk}_{\text{addrA}_\theta}, \text{pk}_{\text{addrB}_\theta}, \text{pk}_{\text{addrC}_\theta}$, for $j \in \{1, \dots, \ell\}$. In addition to the public keys, the j -th instruction of Π is also hard-coded into the script CIScriptInstr_j .

```

1: function  $\text{CIScriptInstr}_j(\sigma_{PV}, pc_\theta, c_{pc_\theta}, \text{insType}_\theta, c_{\text{insType}_\theta},$ 
    $\text{addrA}_\theta, c_{\text{addrA}_\theta}, \text{addrB}_\theta, c_{\text{addrB}_\theta}, \text{addrC}_\theta, c_{\text{addrC}_\theta})$ 
2:    $\text{CheckMSigVerify}_{\text{pk}_{PV}}(\sigma_{PV});$ 
3:    $\text{CheckCommVerify}_{\text{pk}_{pc_\theta}}(pc_\theta, c_{pc_\theta});$ 
4:    $\text{CheckCommVerify}_{\text{pk}_{\text{insType}_\theta}}(\text{insType}_\theta, c_{\text{insType}_\theta});$ 
5:    $\text{CheckCommVerify}_{\text{pk}_{\text{addrA}_\theta}}(\text{addrA}_\theta, c_{\text{addrA}_\theta});$ 
6:    $\text{CheckCommVerify}_{\text{pk}_{\text{addrB}_\theta}}(\text{addrB}_\theta, c_{\text{addrB}_\theta});$ 
7:    $\text{CheckCommVerify}_{\text{pk}_{\text{addrC}_\theta}}(\text{addrC}_\theta, c_{\text{addrC}_\theta});$ 
8:   if  $((pc_\theta = j) \wedge (\text{insType}_j \neq \text{insType}_\theta \vee \text{addrA}_j \neq$ 
    $\text{addrA}_\theta \vee \text{addrB}_j \neq \text{addrB}_\theta \vee \text{addrC}_j \neq \text{addrC}_\theta))$ 
   then
9:     return True;
10:  else
11:    return False.

```

$\text{DisproveProgram} :=$

$$\left(in = [(\text{CommitInstruction}, 0, \text{CIScriptInstr}_{pc_\theta})], \right.$$

$$wit = [(\sigma_{PV}, pc_\theta, c_{pc_\theta}, \text{insType}_\theta, c_{\text{insType}_\theta},$$

$$\text{addrA}_\theta, c_{\text{addrA}_\theta}, \text{addrB}_\theta, c_{\text{addrB}_\theta}, \text{addrC}_\theta, c_{\text{addrC}_\theta})],$$

$$out = [(d\$, \text{CheckSig}_{\text{pk}_V})] \Big). \quad (13)$$

A.1.3. Challenge Read. V starts the challenge by publishing the `ChallengeRead` transaction (cf. Eq. (14)), spending the `CommitInstruction` output¹².

$$\begin{aligned} \text{ChallengeRead} := & \\ (in = & [(CommitInstruction, 0, CheckMSig_{pk_{PV}})], \\ wit = & [(\sigma_{PV})], \\ out = & [(d\mathbb{B}; \langle ReadChallScript_1, TL(\Delta) \wedge CheckSig_{pk_V} \rangle)]). \end{aligned} \quad (14)$$

The script `ReadChallScriptj`, with $j \in \{1, \dots, 5\}$ is defined as follows:

$$\text{ReadChallScript}_j := CheckMSig_{pk_{PV}} \wedge CheckComm_{pk_{Node_{d_{5-j}}}}.$$

The parties engage in the read bisection game (cf. [31, Section B.2]). The game is played over the sequence $\mathcal{P}_R := (MR_N, \dots, M_N[addrA_\theta])$, namely, a path from the root to one of the leaves in $MerkleTree_{M_N}$, i.e., the Merkle tree of the memory at step \mathcal{N} . P responds by publishing the `ReadResponse1` transaction (cf. Eq. (15)), committing to $Node_{d_4} := \mathcal{P}_R[d_4]$ in the witness, where $d_4 = 1 \cdot 2^4$.

$$\begin{aligned} \text{ReadResponse}_1 := & \\ (in = & [(ChallengeRead, 0, ReadChallScript_1)], \\ wit = & [(\sigma_{PV}, Node_{d_4}, c_{Node_{d_4}})], \\ out = & [(d\mathbb{B}; \langle ReadRespScript_1, TL(\Delta) \wedge CheckSig_{pk_P} \rangle)]). \end{aligned} \quad (15)$$

`ReadRespScripti` with $i \in \{1, \dots, 5\}$ is defined as:

$$\text{ReadRespScript}_i := CheckMSig_{pk_{PV}} \wedge CheckComm_{pk_{b'_{5-i}}}.$$

Then, V publishes `ReadChallenge1` transaction (cf. Eq. (16)), committing to bit b'_4 in the witness, where $b'_4 = 1$ if V agrees with $Node_{d_4}$, and $b'_4 = 0$ otherwise.

$$\begin{aligned} \text{ReadChallenge}_1 := & \\ (in = & [(ReadResponse_1, 0, ReadRespScript_1)], \\ wit = & [(\sigma_{PV}, b'_4, c_{b'_4})], \\ out = & [(d\mathbb{B}; \langle ReadChallScript_2, TL(\Delta) \wedge CheckSig_{pk_V} \rangle)]); \end{aligned} \quad (16)$$

P and V continue playing the read bisection game by publishing transactions `ReadResponsei` (cf. Eq. (17)) and `ReadChallengej` (cf. Eq. (18)), respectively, with $i = 2, \dots, 5$ and $j = 1, \dots, 4$.

$$\begin{aligned} \text{ReadResponse}_i := & \\ (in = & [(ReadChallenge_{i-1}, 0, ReadChallScript_i)], \\ wit = & [(\sigma_{PV}, Node_{d_{5-i}}, c_{Node_{d_{5-i}}})], \\ out = & [(d\mathbb{B}; \langle ReadRespScript_i, TL(\Delta) \wedge CheckSig_{pk_P} \rangle)]); \end{aligned} \quad (17)$$

12. We explain how *Challenge Read* works by presenting a challenge to $valA_\theta$; the process for challenging $valB_\theta$ is analogous.

$$\text{ReadChallenge}_j :=$$

$$\begin{aligned} (in = & [(ReadResponse_j, 0, ReadRespScript_j)], \\ wit = & [(\sigma_{PV}, b'_{5-j}, c_{b'_{5-j}})], \\ out = & [(d\mathbb{B}; \langle ReadChallScript_{j+1}, TL(\Delta) \wedge CheckSig_{pk_V} \rangle)]); \end{aligned} \quad (18)$$

$$\text{where } d_{5-i} = 1 \cdot 2^{5-i} + \sum_{k=5-i+1}^4 b'_k \cdot 2^k.$$

Then, V publishes the `ReadChallenge5` transaction (cf. Eq. (19)). In total, V has committed to the bits b'_4, \dots, b'_0 . These bits determine the last element on the path \mathcal{P}_R upon which P and V agree. Let \mathcal{N}_{Mer} be the corresponding integer, computed as $\mathcal{N}_{Mer} = \sum_{k=0}^4 b'_k \cdot 2^k$.

$$\begin{aligned} \text{ReadChallenge}_5 := & \\ (in = & [(ReadResponse_5, 0, ReadRespScript_5)], \\ wit = & [(\sigma_{PV}, b'_0, c_{b'_0})], \\ out = & [(d\mathbb{B}; \langle HashReadScript_1, \dots, HashReadScript_{20}, \\ & RootReadScript_1, \dots, RootReadScript_{32}, ValueAScript, \\ & TL(\Delta) \wedge CheckSig_{pk_V} \rangle)]). \end{aligned} \quad (19)$$

The integer \mathcal{N}_{Mer} , chosen by V , conditions which Tapleaves P can unlock to spend the `ReadChallenge5` output. We can distinguish three cases.

(A) Commit Read. The point of disagreement is between two consecutive elements of the path \mathcal{P}_R , excluding the first and the last. P publishes the `CommitRead1` transaction (cf. Eq. (20)) to spend the `ReadChallenge5` output. To do so, P provides a witness that unlocks one of the scripts `HashReadScript1, ..., HashReadScript20`. Each script hard-codes the public key of a pair of nodes belonging to $\{Node_{d_0}, \dots, Node_{d_4}\}$, the first being the parent node in $MerkleTree_{MR_N}$ and the second being the child node¹³. Additionally, P provide a sibling node `Nsib`, claiming whether it is the left or right sibling by committing to the bit v_{pos} , the \mathcal{N}_{Mer} -th bit of $addrA_\theta$. To unlock the script, it must hold that the child node, when concatenated with the sibling node, hashes to the parent node.

We present the pseudocode of the script in `HashReadScript1` in Algorithm 15. The scripts `HashReadScript2, ..., HashReadScript20` are identical except for the public keys hard-coded to set the parent and the child nodes, and the mapping from \mathcal{N}_{Mer} to the appropriate Tapleaf.

$$\text{CommitRead1} :=$$

$$\begin{aligned} (in = & [(ReadChallenge_5, 0, HashReadScript_i)], \\ wit = & [(\sigma_{PV}, \mathcal{N}_{Mer}, c_{\mathcal{N}_{Mer}}, v_{pos}, c_{v_{pos}}, c_{addrA_\theta}, Nsib, \\ & Npar, c_{Npar}, Nchild, c_{Nchild})], \\ out = & [(d\mathbb{B}; \langle CommitRead1Script, TL(\Delta) \wedge CheckSig_{pk_P} \rangle)]). \end{aligned} \quad (20)$$

13. Since any node can be the parent of any other, we need 20 scripts to capture all the possibilities.

Algorithm 15 The script HashReadScript₁. The bit $v_{pos} \in \{0, 1\}$ represents the position of the child node ($v_{pos} = 0$ means that Node_{d₀} is the left child of Node_{d₄}, $v_{pos} = 1$ means the opposite. Nsib is the sibling node of Node_{d₀} that P presents. In the setup phase, the public keys $pk_{N_{Mer}}$, pk_{addrA_θ} , $pk_{Node_{d_4}}$, $pk_{Node_{d_0}}$, are hard-coded in the script.

```

1: function HashReadScript1( $\sigma_{PV}$ ,  $\mathcal{N}_{Mer}$ ,  $c_{N_{Mer}}$ ,  $v_{pos}$ ,  $c_{v_{pos}}$ ,
    $c_{addrA_\theta}$ , Nsib, Noded4,  $c_{Node_{d_4}}$ , Noded0,  $c_{Node_{d_0}}$ )
2:   CheckMSigVerify $pk_{PV}$ ( $\sigma_{PV}$ );
3:   CheckCommVerify $pk_{N_{Mer}}$ ( $\mathcal{N}_{Mer}$ ,  $c_{N_{Mer}}$ );
4:    $\triangleright$  Since  $V$  committed to  $\mathcal{N}_{Mer}$ ,  $P$  does not know
       $sk_{N_{Mer}}$ . Therefore, to satisfy this guard, has to provide
      the commitment that  $V$  made  $\triangleleft$ 
5:   if CountZeroes( $\mathcal{N}_{Mer}$ )  $\neq 5 - 1$  then
6:      $\triangleright$  Since Noded4 is the parent node here,
        CountZeroes( $\mathcal{N}_{Mer}$ ) should be 4.  $\triangleleft$ 
7:     return False;
8:   CheckCommVerify $pk_{addrA_\theta[N_{Mer}]}$ ( $v_{pos}$ ,  $c_{addrA_\theta[N_{Mer}]}$ );
9:    $\triangleright$  The whole public key  $pk_{addrA_\theta}$  is hard-coded in the
      script, but only the  $\mathcal{N}_{Mer}$ -th entry is used  $\triangleleft$ 
10:  CheckCommVerify $pk_{Node_{d_4}}$ (Noded4,  $c_{Node_{d_4}}$ );  $\triangleright$  Parent
      node
11:  CheckCommVerify $pk_{Node_{d_0}}$ (Noded0,  $c_{Node_{d_0}}$ );  $\triangleright$  Child node
12:  if  $v_{pos} = 0$  then
13:    if  $H(\text{Node}_{d_0} || \text{Nsib}) = \text{Node}_{d_4}$  then
14:      return True;
15:    else
16:      return False;
17:  else
18:    if  $H(\text{Nsib} || \text{Node}_{d_0}) = \text{Node}_{d_4}$  then
19:      return True;
20:    else
21:      return False.

```

V can punish P if they equivocate either on Npar, Nchild, v_{pos} by publishing PunishRead1 (cf. Eq. (21)), which requires to unlock the CommitRead1Script (cf. Algorithm 16) script.

Algorithm 16 The script CommitRead1Script. The public keys pk_{Npar} , pk_{Nchild} , and pk_{addrA_θ} are hard-coded during the setup.

```

1: function CommitRead1Script( $\sigma_{PV}$ ,  $c_0$ ,  $c_1$ )
2:   CheckMSigVerify $pk_{PV}$ ( $\sigma_{PV}$ );
3:   for  $i = 1, \dots, |Npar|_{bit}$  do
4:     if Equivocation( $pk_{Npar[i]}$ ,  $c_0$ ,  $c_1$ ) = True  $\vee$ 
        Equivocation( $pk_{Nchild[i]}$ ,  $c_0$ ,  $c_1$ ) = True  $\vee$ 
        Equivocation( $pk_{addrA_\theta[i]}$ ,  $c_0$ ,  $c_1$ ) = True then
5:       return True;
6:   return False.

```

PunishRead1 :=

$$\left(in = [(CommitRead1, 0, CommitRead1Script)], \quad (21) \right. \\ \left. wit = [(\sigma_{PV}, c_0, c_1)], out = [(d\mathbb{B}; CheckSig_{pk_V})] \right).$$

(B) Commit Value A. If V agrees with every element that P committed (i.e., $b'_4 = \dots = b'_0 = 1$), \mathcal{N}_{Mer} is set to 31. The point of disagreement is between the last interme-

diated node published by P , Node_{d₀}, and $valA_\theta$; To spend the ReadChallenge₅ output, P unlocks ValueAScript. ValueAScript is analogous to HashReadScript_i with the following differences: (i) CountZeroes(\mathcal{N}_{Mer}) = 0; (ii) the parent node is Node_{d₀}; (iii) the child node is not one of the nodes Node_{d₄}, ..., Node_{d₀}, but $valA_\theta$ instead.

P publishes CommitRead2 transaction (analogous to CommitRead1, but unlocking ValueAScript instead). V can publish transaction PunishRead2 (analogous to PunishRead1) if P equivocates on the values committed in the CommitRead2 transaction.

(C) Commit Read Root. If V disagrees with every element that P committed (i.e., $b'_4 = \dots = b'_0 = 0$), \mathcal{N}_{Mer} is set to 0. The point of disagreement is between the last intermediate node published by P , Node_{d₀}, and $MR_{\mathcal{N}}$. P unlocks one of the leaves RootReadScript₁, ..., RootReadScript₃₂, according to which number \mathcal{N} V committed at the end of the dispute bisection game. We provide RootReadScript_i in Algorithm 17.

Algorithm 17 The script RootReadScript_i. In the setup phase, the public keys $pk_{N_{Mer}}$, $pk_{\mathcal{N}}$, $pk_{Node_{d_0}}$, pk_{MR_i} are hard-coded in the script.

```

1: function RootReadScripti( $\sigma_{PV}$ ,  $\mathcal{N}_{Mer}$ ,  $c_{N_{Mer}}$ ,  $v_{pos}$ ,  $c_{v_{pos}}$ ,
    $c_{addrA_\theta}$ , Nsib,  $\mathcal{N}$ ,  $c_{\mathcal{N}}$ , Noded0,  $c_{Node_{d_0}}$ ,  $MR_i$ ,  $c_{MR_i}$ )
2:   CheckMSigVerify $pk_{PV}$ ( $\sigma_{PV}$ );
3:   CheckCommVerify $pk_{N_{Mer}}$ ( $\mathcal{N}_{Mer}$ ,  $c_{N_{Mer}}$ );
4:    $\triangleright$  Since  $V$  committed to  $\mathcal{N}_{Mer}$ ,  $P$  does not know
       $sk_{N_{Mer}}$ . Therefore, to satisfy this guard,  $P$  has to
      provide the commitment that  $V$  made  $\triangleleft$ 
5:   CheckCommVerify $pk_{\mathcal{N}}$ ( $\mathcal{N}$ ,  $c_{\mathcal{N}}$ );
6:    $\triangleright$   $P$  has to provide the commitment that  $V$  made in the
      dispute phase  $\triangleleft$ 
7:   if CountZeroes( $\mathcal{N}$ )  $\neq i$  then
8:     return False;
9:   if CountZeroes( $\mathcal{N}_{Mer}$ )  $\neq 5$  then  $\triangleright$   $\mathcal{N}_{Mer}$  must be equal
      to 0
10:    return False;
11:   CheckCommVerify $pk_{addrA_\theta[N_{Mer}]}$ ( $v_{pos}$ ,  $c_{addrA_\theta[N_{Mer}]}$ );
12:   CheckCommVerify $pk_{Node_{d_0}}$ (Noded0,  $c_{Node_{d_0}}$ );
13:    $\triangleright$  for any RootReadScripti, Noded0 is always the child
      node  $\triangleleft$ 
14:   CheckCommVerify $pk_{MR_i}$ (Noded0,  $c_{Node_{d_0}}$ );
15:   if  $v_{pos} = 0$  then
16:     if  $H(\text{Node}_{d_0} || \text{Nsib}) = MR_i$  then
17:       return True;
18:     else
19:       return False;
20:   else
21:     if  $H(\text{Nsib} || \text{Node}_{d_0}) = MR_i$  then
22:       return True;
23:     else
24:       return False.

```

P unlocks RootReadScript_i by publishing the CommitRead3 transaction (cf Eq. (22)).

CommitRead3 :=

$$\left(in = [(\text{ReadChallenge}_5, 0, \text{ReadRootScript}_i)], \right.$$

$$wit = [(\sigma_{PV}, \mathcal{N}_{Mer}, c_{\mathcal{N}_{Mer}}, \mathcal{N}, c_{\mathcal{N}}, \text{Node}_{d_0}, c_{\text{Node}_{d_0}})],$$

$$\left. out = [(d\mathbb{B}; \langle \text{CommitRead3Script}, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_P} \rangle)] \right).$$
(22)

V can punish P if they equivocate on Node_{d_0} , MR_i or $\text{addr}A_\theta$ by publishing PunishRead3 (cf. Eq. (23)), which unlocks CommitRead3Script , analogous to CommitRead1Script but with pk_{MR_i} , $\text{pk}_{\text{Node}_{d_0}}$ instead of pk_{Npar} , pk_{Nchild} .

PunishRead3 :=

$$\left(in = [(\text{CommitRead3}, 0, \text{CommitRead3Script})], \right.$$
(23)

$$wit = [(\sigma_{PV}, c_0, c_1)], out = [(d\mathbb{B}; \text{CheckSig}_{\text{pk}_V})].$$

A.1.4. Challenge Write. V challenges the result of the writing operation. Specifically, V claims that P is writing $\text{val}C'_\theta \neq \text{val}C_\theta$ in $M_{N'}[\text{addr}C_\theta]$ in their local VM execution¹⁴. As a result, the memory root $MR_{N'}$ is incorrect.

The parties engage in the write bisection game (cf. [31, Section B.3]) over the sequences $\mathcal{P}_W := (MR_N, \dots, M_N[\text{addr}C_\theta])$ and $\mathcal{P}'_W := (MR_{N'}, \dots, M'_{N'}[\text{addr}C_\theta])$, that are paths in the merkle trees MerkleTree_{M_N} and $\text{MerkleTree}_{M_{N'}}$, respectively. The transactions and locking scripts in the challenge write branch of the protocol closely follow the structure of those in the challenge read branch, with the following differences:

- The structure of the WriteResponse_i transaction is analogous to ReadResponse_i transaction but, in the witness, P provides two values (and their commitments) instead of one. These values are the d_{5-i} -th elements of \mathcal{P}_W and \mathcal{P}'_W , respectively.
- As long as V agrees on the elements of the path \mathcal{P}_W , they focus on finding the disagreement in the path \mathcal{P}'_W . In the WriteChallenge_j transaction (analogously to ReadChallenge_j), V sets (and commits to) the bit $b'_{5-j} = 0$ if V agrees with the element of \mathcal{P}'_W provided by P . Otherwise, V sets (and commits to) the bit $b'_{5-j} = 1$. However, once V finds a disagreement in an element of \mathcal{P}_W , from that point on, V focuses on \mathcal{P}_W and set the bit b'_{5-j} as in the *Challenge Read* branch.

During the write bisection game, P commits to the pairs nodes $\{(\text{Node}_{d_4}, \text{Node}'_{d_4}), \dots, (\text{Node}_{d_0}, \text{Node}'_{d_0})\}$, where $\text{Node}_{d_4}, \dots, \text{Node}_{d_0} \in \mathcal{P}_W$ and $\text{Node}'_{d_4}, \dots, \text{Node}'_{d_0} \in \mathcal{P}'_W$. Analogous to the *Challenge Read* branch, V commits bit by bit to an integer $\mathcal{N}_{Mer} = \sum_{k=0}^4 b'_k \cdot 2^k$, which conditions how P can unlock WriteChallenge_5 . There are three cases.

Note that P does not explicitly know which pair of elements in \mathcal{P}_W or \mathcal{P}'_W V disagrees with. However, as long as

¹⁴ We assume P commits correctly to $\text{val}C_\theta$ in the witness of the CommitInstruction transaction, regardless of local execution. For example, if $\text{insType}_\theta := \text{ADD}$, then $\text{val}A_\theta + \text{val}B_\theta = \text{val}C_\theta$. If $\text{val}C_\theta$ is incorrect, V can challenge $\text{val}A_\theta$ or $\text{val}B_\theta$.

P is able to provide a pair of nodes ($Npar, Nchild$) for \mathcal{P}_W , a pair of nodes ($Npar', Nchild'$) for \mathcal{P}'_W , and a node $Nsib$ such that $H(Nsib || Nchild) = Npar$ and $H(Nsib || Nchild') = Npar'$, they will be able to unlock WriteChallenge_5 .

(A) CommitWrite. The point of disagreement is between two consecutive elements of \mathcal{P}_W or between two consecutive elements of \mathcal{P}'_W , excluding for both paths the first and the last elements. P can unlock one of the scripts $\text{HashWriteScript}_1, \dots, \text{HashWriteScript}_{20}$ via publishing the CommitWrite1 transaction (cf. Eq. (24)). Each script HashWriteScript_i is identical to script HashReadScript_i , for $i = 1, \dots, 20$, except that it also verifies the commitments of the parent and child nodes on the path \mathcal{P}'_W (their public key are hard-coded in the script accordingly) and it checks whether they really are a parent-child pair. We present script HashWriteScript_1 in the extended version of this paper [31].

CommitWrite1 :=

$$\left(in = [(\text{WriteChallenge}_5, 0, \text{HashWriteScript}_i)], \right.$$

$$wit = [(\sigma_{PV}, \mathcal{N}_{Mer}, c_{\mathcal{N}_{Mer}}, v_{pos}, c_{v_{pos}}, c_{\text{addr}C_\theta}, Nsib,$$

$$Npar, c_{Npar}, Nchild, c_{Nchild}, Npar', c_{Npar'}, Nchild',$$

$$c_{Nchild'})],$$

$$\left. out = [(d\mathbb{B}; \langle \text{CommitWrite1Script}, \text{TL}(\Delta) \wedge \text{CheckSig}_{\text{pk}_P} \rangle)] \right);$$
(24)

The script $\text{CommitWrite1Script}$ is identical to CommitRead1Script (cf. Algorithm 16) except that it also checks for potential equivocation on $Npar'$, $Nchild'$, and $\text{addr}C_\theta$ rather than $\text{addr}A_\theta$. As a consequence, the PunishWrite1 transaction is analogous to PunishRead1 . Thus, if P equivocates while committing to $Npar, Nchild, Npar', Nchild'$, V can claim all the coins locked in the multisignature.

(B) Commit Value C. $\mathcal{N}_{Mer} = 31$, the point of disagreement is between Node_{d_0} and $\text{val}C_\theta$ or between Node'_{d_0} and $\text{val}C_\theta$. This case is analogous to the “commit value A” case of the *Challenge Read* branch. For ValueCScript , the difference with HashWriteScript_i is that: (i) $\text{CountZero} = 0$; (ii) the parent nodes are Node_{d_0} and Node'_{d_0} , and (iii) the child node is $\text{val}C_\theta$.

P publishes CommitWrite2 transaction (analogous to CommitWrite1 , but unlocking ValueCScript instead). V can publish transaction PunishWrite2 (analogous to PunishWrite1) if P equivocates on the values committed in the CommitWrite2 transaction.

(C) Commit Write Root. $\mathcal{N}_{Mer} = 0$, the point of disagreement is between MR_N and Node_{d_0} or between $MR_{N'}$ and Node'_{d_0} . This case is analogous to the “commit read root” case of the *Challenge Read* branch. The script RootWriteScript_i , with $i \in \{0, \dots, 31\}$ is the same as RootReadScript_i but takes as additional inputs Node'_{d_0} , MR_{i+1} (their public key are hard-coded in the script accordingly), and takes as input $c_{\text{addr}C_\theta}$ instead of $c_{\text{addr}A_\theta}$. We present script RootWriteScript_i in the extended version of this paper [31].